

A Closer Look at the Gambler's Fallacy*

Lucas de Lara

Department of Economics, Binghamton University - State University of New York
lpanicodelar@binghamton.edu

February 2, 2026

Abstract

A classic explanation for the gambler's fallacy is that subjects believe that sequences of tosses from a fair coin should be representative of the randomness of the uniform distribution, and so should not have observable patterns. I introduce an information-theoretic formalization of this representativeness heuristic in terms of complexity and contrast it to the existing recency-weighted reversal model by Rabin and Vayanos. In order to test between these explanations, I collect rich choice and belief data from subjects predicting the next item from fully randomized sequences of binary outcomes, allowing me to take the analysis to the level of individual sequences. The basic results confirm the existence of the gambler's fallacy in the aggregate. However, there is also significant heterogeneity among subjects. I identify four types, depending on whether they report correct beliefs or incorrect beliefs that go in the gambler's fallacy direction, its opposite or a mix of both. Taking this heterogeneity into account, both models perform well when looking at an aggregate level, but a closer look at individual sequences reveals violations of the representativeness model which lead to a superior performance of the recency-weighted reversal model. The main component of this superior performance comes from the recency bias that subjects exhibit, which is not accounted for in the representativeness model. I then extend the theoretical and experimental paradigm to predictions about the next two outcomes. The results largely reiterate the conclusions from the original experiment.

JEL Classifications:

Keywords: Gambler's Fallacy, Representativeness, Complexity, Empirical Entropy

*Click here for lastest version. I thank the members of the Cognition and Decision Laboratory at Columbia University and I recognize the financial support provided by the Columbia Experimental Laboratory for Social Sciences (CELSS).

1 Introduction

Imagine yourself observing someone toss a fair coin ten times in front of you, resulting in the sequence: head, head, head, head, head, head, head, head, head, head. This outcome was just as statistically likely as any of the other ones that could have happened, but a substantial and long-standing body of literature indicates that such an occurrence challenges commonly held but incorrect intuitions about randomness. The *gambler’s fallacy* refers to a bias from which individuals erroneously anticipate too many reversals in sequences of independent random events. In the case mentioned, some might believe that, if the coin is truly fair and were to be tossed again, the next outcome would be more likely to – or even would *have to* – result in a tail.

The gambler’s fallacy has long received learned attention. Oskarsson et al. (2009) provide a detailed survey of the interest and investigation it has received from several fields. Documentation of the phenomenon goes back at least as far back as the late eighteenth century, with Laplace noting that people would bet on numbers that they expected to be “due” for a draw in the French lottery, believing that “because the number has not been drawn for a long time, it, rather than the others, ought to be drawn on the next draw.” (Dale and Laplace 1825/1995, p.92) His observations highlight a persistent tendency to misinterpret independent random events, expecting outcomes to self-correct in the short term to align with long-term statistical properties. This expectation of self-correction can even apply to one’s own behavior. In the Zenith radio experiments in telepathy, happening in the late 1930s, listeners were invited to send their predictions over sequences of binary outcomes. Skinner (1942) was the first to point out that people’s guesses were sequentially autocorrelated and exhibited what was to be known as the gambler’s fallacy effect, first bringing it to the attention of psychologists. Since then, a large literature has developed in pursuit of explanations for how and why people hold these erroneous beliefs.

Despite the substantial size of this literature, existing theories remain lacking in either formalization or complete empirical testing. In this paper, I formalize a classical theory of the gambler’s fallacy and compare it to the most recent model in the literature. In order to do so, I conducted an experiment that provided the data

richness required to properly compare these two theories. Kahneman and Tversky (1972) attributed the gambler’s fallacy to the representativeness heuristic. I return to a close reading of their description of this application, which reveals a connection between representativeness in the case of a fair coin and the complexity of the sequences of outcomes that it generates. I then develop a formalization which closely follows their words. This formalization brings tools from information theory to characterize randomness as information complexity, measured by the (higher order) empirical entropy of sequences. This measure includes moments of higher order than the basic proportions that have been the focus of existing models. In the most recent of such models, Rabin and Vayanos (2010) model the gambler’s fallacy using two parameters, a reversal parameter which captures the proportion in previous outcomes, and a discount parameter which captures a recency bias effect. This model, however, still lacks a detailed empirical analysis, especially one that can differentiate between the two effects created by the two parameters.

Conducting the appropriate empirical analysis of these two theoretical explanations and comparing their performance requires a richness of data that has been missing in the literature thus far. Previous experiments have often used tasks that are difficult to incentivize or analyze, have often narrowly focused on a limited or pre-selected set of sequences of outcomes and have utilized limited data collection methods. By running an online experiment in which I collect, from a broader population, comprehensive choice and belief data for all possible sequences, under randomization, I obtain a richer dataset that enables detailed sequence-by-sequence analysis which allows for differentiating between competing explanations. The data confirms classical results regarding the choice of the next outcome when faced with full runs of the same outcome. However, closer inspection reveals important heterogeneities across subjects. Some subjects behave rationally, giving the correct answer, while a similar number indeed follows the gambler’s fallacy, and yet others, in smaller numbers, actually behave in the opposite way. Focusing on the appropriate types of subjects, both the Rabin and Vayanos model and the entropy analysis perform well at the aggregate level, with highly significant estimated parameters. Taking the analysis to

the level of individual sequences reveals the differences between the two: the Rabin and Vayanos model captures very well the observations in the data, while for some sequences, subjects' behavior contradicts the predictions coming from entropy.

The conceptual move to make the gambler's fallacy a case of a broader feature of cognition was made by Kahneman and Tversky (1972). They considered the gambler's fallacy to be a specific instantiation of the representativeness heuristic, wherein individuals expect even small samples to reflect the characteristics of the larger population or process from which they are drawn.¹ When applying the representativeness heuristic to the case of the gambler's fallacy, Kahneman and Tversky included in their definition of representativeness a component relating to people's intuitions about randomness, which is summed as a lack of regularity, a component that has escaped a proper formalization in existing models. Based on this notion of randomness, they proceed to say that, in the case of outcomes from a fair coin, they "venture that only HTTHHTH appears really random. For four tosses, there may not be any" (*ibid.*, p. 436). The cautious, non-committal language, though, highlights the likely absence of any formal and specific criteria for such judgements. Nonetheless, as I will show, a close reading of their verbal descriptions reveals links between their interpretation of randomness and notions of data compression, familiar from Information Theory Shannon (1948). Apparent randomness is linked to complexity, an absence of patterns

¹Since the seminal paper, the representativeness heuristic has gone on to produce an expansive literature of its own. See, for example, Tversky and Kahneman (1974), Tversky and Kahneman (1982), Bar-Hillel (1984). For the purposes of this paper, however, the most important points to notice regarding this literature are the following. The main application of the heuristic has been to judgements of categorization (how likely an object is to belong to a certain set/category). The application to the gambler's fallacy has received comparatively little attention (it receives, e.g., passing mention in Bar-Hillel (1980, p. 588) and Bar-Hillel (1984, p. 103)). One kind of formalization that has existed in the literature is regarding base-rate neglect in Bayesian inference (e.g. Grether 1980), but this cannot explain the gambler's fallacy as it goes in the opposite direction of inference by evidence. The closest formalization to the one being presented here comes from Tenenbaum and Griffiths (2001), which formalizes representativeness as relative model evidence in the context of Bayesian inference. This approach has some important similarities with my formalization, to be discussed when it is introduced in subsection 2.2, but will also be shown to have three important differences: (i) it is not a model of belief formation and prediction, and is more appropriate to what I call in section 3 "randomness judgement tasks", which is indeed the type of task that their experiment on coin flips is; (ii) it lacks the connection to information theory and complexity that I establish; (iii) it leaves open what is the appropriate space of alternative hypotheses/models, whereas it will be shown that the complexity connection automatically establishes a relevant space of models.

in observed outcomes, which represents incompressibility of information.

I then follow this interpretation and formalize the gambler's fallacy as an expectation that outcomes from a fair coin will maintain a high level of information complexity in the sequence of outcomes that they produce. This complexity is measured by empirical entropy (Manzini 2001; Ferragina and Manzini 2005), which measures the proportions of outcomes coming after various combinations of previous outcomes. This means that sequences with low empirical entropy are those with clearer patterns, which allow for greater information compression, while high empirical entropy implies a lack of such patterns, lower availability for compression, and therefore high complexity. The connection to randomness and representativeness of the uniform distribution of a fair coin comes from the equivalence of empirical entropy and likelihood inference within the space of Markov chains (Gagie 2006). Sequences with high complexity are those for which the best model fit comes closer to the uniform distribution, which is the case of maximum entropy.

I contrast this formalized information theoretic version of the representativeness heuristic to Rabin (2002) and Rabin and Vayanos (2010). In the first one, the gambler's fallacy is modeled as a belief that outcomes are being drawn without replacement from a finite number of possible outcomes. The second, and more recent one, is a recency-weighted reversal model in which the gambler's fallacy is the joint result of two parameters: a reversal parameter, which captures the first moment of proportions and the expectation that this proportion will tend towards a balanced one, and a discount factor, which captures a recency bias in that more recent outcomes are weighted more heavily than those further past.

The experimental literature, in turn, has focused on the common paradigmatic tasks of choosing future outcomes after seeing a given sequence of outcomes, constructing random sequences and judging the randomness of given sequences. A common issue shared by all these tasks, similarly observed by Rabin (2002, p. 782) and repeated by Rabin and Vayanos (2010, fn. 10 and 11), is that, beyond just often not having been incentivized in existing studies, they are hard to incentivize at all. For all of them, there are no right or wrong answers, and it is hard to argue conclusively

that there is anything wrong with any behavior observed from subjects. This is so even if specific and suggestive patterns are found in the data. The most natural solution is to focus on beliefs. In the case of a fair coin, the only correct belief to hold over the next outcome is the uniform Bernoulli distribution, which matches the true known data generating process. I leverage this fact to design a properly incentivized task. Nonetheless, it's still useful to collect choice data as well to keep within the tradition of the literature, and so I collect both choice and belief data, in this order and required to be consistent, for each sequence.

In contrast to most of the literature, Benjamin et al. (2017) made significant advances with a series of well-incentivized experiments. However, a significant limitation in their approach is that they only asked for belief responses for sequences of full runs of a single outcome. For other sequences, choice data was collected. This means that they could estimate the Rabin and Vayanos model with sequences of the former type, but could not precisely test it against those of the latter type. Moreover, they could not differentiate between the two effects present in the model, the reversal effect and the recency effect. This paper, on the other hand, allows for a detailed test of the model, and results show that the recency effect, in particular, plays a major role in explaining the fitting of the model to the observed data.

To obtain the detailed data necessary for the analysis mentioned above, I designed and ran an online experiment which proceeded as follows. Subjects were faced with a series of sequences of outcomes from flipping a computerized fair coin, heads and tails, and were asked to make two responses on each sequence. First, they had to choose by clicking on which outcome, head or tail, they wanted to bet on being the next one to happen. Second, they had to move sliders to give, in integer numbers from 0 to 100, their probability prediction for the likelihood of the next outcome being either a head or a tail. The sliders were programmed so that their probability response had to be coherent with their previous choice response, that is, if the subject chose head as the next outcome, their probability response for head had to be at least 50. This helps linking choice and probability responses as well as making choice decisions more meaningful for subjects which really believe that the probabilities are not 50-50.

297 subjects were recruited on the Prolific platform. Each subject faced 54 sequences. The first 50 were fully randomized, and consisted of 2 sequences of length 0, 4 sequences of length 2, 12 sequences of length 4, 20 sequences of length 6, and 12 sequences of length 8. The final 4 rounds consisted, in random order, of 1 sequence of each non-zero length but of the full run type. At the end, subjects answered a final non-incentivized follow-up multiple choice question about their reasoning during the experiment.

Plotting choice and probability response data for the full run sequences confirms a gambler's fallacy effect, especially on choice data as has been classically done in the literature. But a weaker effect on probability as well as a weaker monotonicity on the length of the run than might be expected suggest a more complicated picture. The first major result that departs from the previous literature is the identification of significant heterogeneity among subject types. Based on their probability responses to the last 4 rounds, subjects were classified into four types:

1. *rational*: Subjects who responded 50-50 probabilities in all rounds. Their responses demonstrate an understanding of the behavior of a fair coin and probabilistic independence.
2. *gambler's fallacy*: Subjects whose responses were consistent with the gambler's fallacy, expecting reversals of runs to be more likely.
3. *hot hands*: Subjects who showed the opposite pattern to the above, expecting runs to be more likely to continue than to break.
4. *both/mixed*: Subjects showing inconsistent patterns, sometimes predicting reversals and other times continuations.

The respective proportions of these types were 29.62%, 28.61%, 19.52% and 22.22%. Comparing subjects' types to their responses in the follow-up question revealed significant consistencies between the two. 'rational' subjects overwhelmingly respond that past flips do not affect future flips. 'gambler's fallacy' subjects overwhelmingly gave the expected response that the coin should generate balanced

outcomes and therefore reversals were more likely. ‘hot hands’ subjects had more diversity in their responses, but they were consistent in giving responses that logically explain their answers. Then, looking separately by type at probability and choices responses for full runs confirms the patterns that one might expect from this classification, with significant effects in the appropriate direction for ‘gambler’s fallacy’ and ‘hot hands’ subjects as well as monotonicity.

Given subject heterogeneity, the proceeding analysis was done separately by subject type, with a special focus on ‘gambler’s fallacy’ subjects. The first step is to fit both the Rabin and Vayanos model and the empirical entropies model to probability responses for all sequences. This gives highly significant parameters in both cases. In the Rabin and Vayanos case, both reversal and recency parameters are highly significant and of sizes that produce substantial effects on both fronts. For empirical entropies, the parameters for zeroth, first and second order entropies are all significant, but their parameter sizes are decreasing, showing that more complex patterns have less impact. Both perspectives, therefore, confirm at an aggregate level the impact of sequence properties beyond the basic proportion between the two outcomes. But a closer analysis is required to differentiate between the two.

The detailed nature of the data collected is then leveraged to properly compare the two models. The analysis is taken to the level of individual sequences of length up to 6 to contrast mean predictions against the predictions of each model. This level of analysis reveals that the Rabin and Vayanos model performs significantly better than predictions coming from sequence complexity. In particular, there are sequences, most notably the alternating sequences like *head-tail-head-tail...* for which subjects predict a continuation of the pattern rather than a break. This goes against the representativeness model, as a pattern break increases complexity. It is aligned, however, to the recency effect in Rabin and Vayanos model, as one of the outcomes is further to the front of the sequence than the other, and is thus more heavily weighted in the expectation of reversal.

The analysis thus far finds that the Rabin and Vayanos model performs significantly well against the data. The final step of the analysis is to consider whether this

model, which consists of two parameters can improved upon by a generalized model with eight parameters, one for each possible position in the sequences. The results reveal that the estimated generalized parameters are very close to those implied by the two parameters in the original model, confirming once again its excellent fit to the data.

One aspect to be highlighted is that the gambler’s fallacy is particularly puzzling from a bounded rationality perspective. Especially in the context of this paper’s experiment, the task at hand is simple and it’s hard to see what cognitive constraints might be operant and leading to the wrong answers. If subjects have no sense of what outcomes from a fair coin are supposed to be, it would make sense for their answers to follow the pattern of statistical inference, expecting, for example, heads to be more likely if that’s what they observe. This is indeed how the ‘hot hands’ subjects mentioned above behave. But the gambler’s fallacy goes in the opposite direction. By studying this kind of behavior, which might be described as a kind of anti-inference, it might be possible to have a better understanding of the mechanisms of belief formation in general. For belief formation, and predictions generated by them, permeate several, if not all, aspects of economic decision making. This is particularly acute for domains where probabilistic thinking is most essential for reaching conclusions, such as financial decision making.

This paper is structured as follows. Section 2 provides some notation and the theoretical exposition of the two models being considered and compared. Section 3 introduces the experimental design and explains how it was implemented. And section 4 provides the results from analyzing the data collected from the experiment. Section 5 then extends the framework of the previous section into the case of predicting the next two outcomes of the process. A conclusion discusses the results.

2 Theory

2.1 Notation

An agent faces finite binary strings of varying length $n \in \mathbb{N}$. They are denoted $s^n \in 2^\omega$. For example, $s^4 = 0011$ or $s^6 = 101110$. The data generating process producing these strings is given by independent draws from a distribution $p \in \Delta(\{0, 1\})$, and the agent is informed of this process. In the case of a fair coin, p is the uniform Bernoulli distribution. After seeing a sequence s^n , the agent reports a belief q over the next outcome. That is, the agent's responses follow a belief mapping $s^n \mapsto q \in \Delta(\{0, 1\})$. Let q_0 be the probability given for outcome 0. In some cases, the affine transformation $x \mapsto 2x - 1$ and its inverse are used, so that it may be that $q_0 \in [0, 1]$ or $q_0 \in [-1, 1]$.

In the gambler's fallacy, q deviates from p in a systematic way. For example, despite p being uniform, the agent might report $q_0(s^6) < 0.5$ if $s^6 = 010000$.

2.2 Complexity as Representation of Randomness

The gambler's fallacy can be modeled as people expecting that outcomes from a fair coin should not produce noticeable patterns. Instead, they expect that the sequence of outcomes is supposed to maintain a high degree of complexity, which will be shown to mean an absence of patterns of any kind. This complexity is linked to the kind of randomness that is expected from the uniform distribution. In this sense, sequences of higher complexity are more representative of this kind of randomness than patterned sequences. This subsection shows how this idea can be expressed formally via the complexity measure of empirical entropy in a way that captures the original notion of the representativeness heuristic.

Kahneman and Tversky (1972) introduced the idea of the representativeness heuristic and presented the gambler's fallacy as one of its applications. According to this notion of representativeness, the “subjective probability of an event, or sample, is determined by the degree to which it: (i) is similar in essential characteristics to its parent population; and (ii) reflects the salient feature of the process by which it is generated.” In the case of a fair coin, the process is the binary uniform distribu-

tion, linking a representative sample to the kind of randomness that is expected from such a distribution. “A representative sample, then, is similar to the population in essential characteristics, and reflects randomness as people see it; that is, all its parts are representative and none is too regular.” (ibid. p. 436)

A formalization of this idea requires further clarification regarding what it means for a sequence to appear random, and for its parts to not be ‘too regular’. The following words, worth quoting in full, offer strong hints of how one should proceed:

Random-appearing sequences are those whose verbal description is longest.

Imagine yourself dictating a long sequence of binary symbols, say heads and tails. You will undoubtedly use shortcut expressions such as ‘four Ts,’ or ‘H-T three times.’ A sequence with many long runs allows shortcuts of the first type. A sequence with numerous short runs calls for shortcuts of the second type. *The run structure of a random-appearing sequence minimizes the availability of these shortcuts, and hence defies economical descriptions. Apparent randomness, therefore, is a form of complexity of structure.* (ibid. pp. 436–437, both italics and bold mine)

Based on this, I turn to Information Theory to formalize random-appearing sequences as those with a high degree of complexity, which, equivalently, are those with low levels of information compressibility. Inversely, less complex sequences are those with patterns that can be exploited for information compression. This is captured via the complexity measure of (higher order) empirical entropy (Manzini 2001; Ferragina and Manzini 2005).

Let $\Sigma = \{\alpha_1, \dots, \alpha_h\}$ be a finite alphabet. Let s^n be a string of n symbols from Σ , and let n_i , $i = 1, \dots, h$, be the number of occurrences of symbol α_i in s^n . The (zero-th order) empirical entropy of s^n is given by²

$$H_0(s^n) = - \sum_{i=1}^h \frac{n_i}{n} \log \left(\frac{n_i}{n} \right).$$

²The logarithm being base 2 and assuming $0 \log 0 = 0$.

This is equivalent to the entropy of the empirical distribution of the sequence. $nH_0(s^n)$ measures the maximum compression of the sequence that's achievable via a uniquely decodable code with a fixed codeword for each symbol. Note that in the binary case, this simply registers the relative proportion of the two possible symbols. In the case of a fair coin, the relative proportion between heads and tails. It equals 0 if the outcomes are all heads or all tails, and 1 if they're exactly balanced. Therefore, at this zero-th order level of consideration, an expectation of high complexity simply means an expectation that 0s and 1s ought to stay in a balanced proportion.

To capture more complicated patterns that might exist in the sequence or might be expected as more outcomes happen, one needs to consider the possibility of using codewords that depend not only on the symbol itself, but also on the symbols preceding it. For each $w \in \Sigma^k = \Sigma \times \Sigma \times \dots \times \Sigma$ and $\alpha_i \in \Sigma$, let $n_{w\alpha_i}$ denote the number of times that the substring $w\alpha_i$ appears in the string s^n , and let $n_w = \sum_i^h n_{w\alpha_i}$. The k -th order empirical entropy of s^n is given by

$$H_k(s^n) = \sum_{w \in \Sigma^k} \frac{n_w}{n} \left[- \sum_{i=1}^h \frac{n_{w\alpha_i}}{n_w} \log \left(\frac{n_{w\alpha_i}}{n_w} \right) \right],$$

with $H_{k+1}(s) \leq H_k(s)$ for any s and k . Understanding the formula is easier after knowing the procedure for calculating it. For each $w \in \Sigma^k$, one constructs the subsequence of symbols from s^n that come after w , then calculates the zero-th order empirical entropy of this subsequence (this is the value in the brackets), which is then weighted by n_w/n , and then moves on to the next element in Σ^k .

In the binary case, $0 \leq H_k(s) \leq 1$ for any k and s . Sequences with higher empirical entropy, closer to 1, are those to be considered more complex, up to the level of the order being considered. To quickly see how these higher orders differentiate the complexity of sequences with regard to patterns that go beyond basic proportions, consider the example of two sequences of length 8,

$$s^8 = 01010101 \text{ and } \tilde{s}^8 = 10100011.$$

Both have both outcomes in equal proportion, but the second one appears more

complex than the first one, which actually seems highly patterned. And, indeed, we have

$$H_0(s^8) = H_0(\tilde{s}^8) = 1, \text{ but } H_1(s^8) = 0 \text{ while } H_1(\tilde{s}_8^2) \approx 0.8443.$$

While empirical entropy allows for the measuring of complexity of sequences, it's not yet clear that it provides any more rigorous justification for the notion that more complex sequences are those that appear more random, and thus more representative of the uniform distribution. The missing link is provided by the following result.

Theorem 1 (Gagie, 2006).

$$H_k(s^n) = \frac{1}{n} \min_{P \in \mathcal{P}_k} \log \frac{1}{P(s^n)}$$

where \mathcal{P}_k is the set of k -th order Markov chains.

The theorem states the empirical entropy of order k of a sequence s^n , as calculated by the formulas above, coincides with the (averaged by size) log-likelihood of the best explanation for the sequence within the space of k -th order Markov chains.

This shows a direct connection between the complexity of a sequence, measured by empirical entropy, and the best likelihood inference that its outcomes suggest within the space of Markov chains. Notice that the uniform distribution is always in the set \mathcal{P}_k . In the binary case, plugging in the uniform Bernoulli results, for any s^n , in $\frac{1}{n} \log \frac{1}{P(s^n)} = 1$, which is the level of maximum entropy. High complexity, then, is a worst case scenario in likelihood inference, in which no better explanation than the uniform distribution can be suggested by the outcomes of the sequence. This formalizes representativeness in terms of statistical inference while also connecting it to the idea of high complexity in the case of the uniform distribution.

The gambler's fallacy can then be modeled as an expectation that the next outcome in the sequence should be the one that makes the sequence more complex than the other. Given the development in the previous paragraphs, this means that the outcomes also make the sequence more representative of the uniform distribution. To capture the difference that one outcome makes relative to the next, the following variables are introduced:

$$\Delta_0 H_k(s^n) = H_k(s^n0) - H_k(s^n1),$$

in which s^n0 is the sequence s^n with a 0 added at the end, and likewise for s^n1 . The agent's belief q_0 over the next outcome being 0 becomes a function of these differences in empirical entropies $\Delta_0 H_k$.

The following table shows the example of the sequence $s^6 = 000111$, for which $\Delta_0 H_0(s^6) = 0$, $\Delta_0 H_1(s^6) > 0$ and $\Delta_0 H_2(s^6) > 0$.

Sequence	Future Sequence	H_0	H_1	H_2
000111	0001110	0.985228	0.787111	0.571429
	0001111	0.985228	0.393555	0.285714

I take an agnostic position with regards to the precise functional form that relates these variables $\Delta_0 H_k$ to beliefs. Given the theory that the gambler's fallacy consists in believing that complexity-increasing outcomes should be more likely, this means that beliefs will follow a function

$$q_0(s^n) = q_0(\Delta_0 H_0(s^n), \Delta_0 H_1(s^n), \Delta_0 H_2(s^n))$$

which is increasing in its arguments, and such that $q_0(0, 0, 0) = 0.5$.³ In the example above of $s^6 = 000111$, the prediction is that $q_0(000111) > 0.5$.⁴

³Given the binary alphabet being worked with, and that in the experiment sequences will be of length up to 8, only empirical entropies of up to order 2 will be considered, following the results from Gagie (2006)

⁴As discussed in footnote 1 in the introduction, this approach to representativeness has some similarities but mostly differences to that seen in Tenenbaum and Griffiths (2001). They present a measure of representativeness in terms of relative model/hypothesis evidence in terms of Bayesian inference. How representative observed data d is of a model/hypothesis $h_i \in \mathcal{H}$ is given by

$$R(d, h_i) = \log \frac{P(d|h_i)}{\sum_{h_j \in \mathcal{H}} P(d|h_j)P(h_j|\bar{h}_i)}$$

in which $P(h_j|\bar{h}_i)$ is the prior probability of h_j given that h_i is not the correct model/hypothesis. This approach is similar to mine in considering representativeness as *relative* representativeness, comparing likelihoods between different possible models. The first immediate difference is the Bayesian framework with the use of prior beliefs. The second is a lack of connection to information theory and complexity. The third is that this is not extended to belief formation and prediction, and their

Below is another example, this time the case of $s^6 = 010101$, which has $\Delta_0 H_0(s^6) = 0$, $\Delta_0 H_1(s^6) < 0$ and $\Delta_0 H_2(s^6) < 0$. In this case, the prediction is that $q_0(010101) < 0.5$.

Sequence	Future Sequence	H_0	H_1	H_2
010101	0101010	0.985228	0.000000	0.000000
	0101011	0.985228	0.393555	0.393555

2.3 Reversal and Recency

The formalization of the representativeness heuristic in the previous subsection was capable of capturing the effects of higher order complexity. This subsection, in turn, presents the most recent model of the gambler’s fallacy in the literature, the recency-weighted reversal model presented by Rabin and Vayanos (2010). Their model is also capable of capturing the effects on beliefs of more complicated patterns than simple proportions. It models the gambler’s fallacy through the combination of two effects, each with its own parameter, a reversal effect and a recency effect. It’s the latter one that is necessary for creating a gambler’s fallacy effect that goes beyond the first moment of relative numbers of the two outcomes. This model makes predictions that sometimes go in the same and sometimes in the opposite direction of the predictions coming from the representativeness model, opening the way for differentiating between them using experimental data.

In the general version of the model, an agent needs to make predictions on future outcomes based on inferences they make about an underlying moving state. The agent receives noisy signals about the state, and the gambler’s fallacy is modeled as the ‘mistaken belief’ that the noise shocks in the signals are not i.i.d. as they actually are, but rather must exhibit reversals with regard to previous outcomes, which are

experiment indeed only asked subjects to judge how likely four preselected sequences were to be coming from a few alternative coins (fair, alternating, biased and deterministic). The fourth is that the space of possible models \mathcal{H} being considered is left undetermined, whereas the complexity connection establishes a specific set of possible models, namely the set of Markov chains of a specific order. The fifth is that my approach compares the reference distribution, the uniform Bernoulli, to the *best* alternative explanation within this space of possible models.

weighted by a discount factor.

In the case of an agent observing the results from tossing a fair coin, the underlying state can be taken to be fixed, or non-existent, leading to a belief function that takes the form⁵

$$q_0(s^n) = \alpha \sum_{t=0}^{\infty} \delta^{t+1} s_{n-t}^n \in [-1, 1]$$

where q_0 is probability of 0, and s_{n-t}^n codes the value in position $n - t$ of the sequence:

$$s_{n-t}^n = \begin{cases} 1, & \text{if } = 1, \\ 0, & \text{if } = \emptyset, \\ -1, & \text{if } = 0. \end{cases}$$

The parameters $(\alpha, \delta) \in [0, 1]^2$ capture, respectively, the reversal and recency effects. The reversal effect is stronger the closer α is to 1, while the recency effect is stronger the closer δ is to 0.

Notice that if $\delta = 1$, that is, if there's no recency effect, beliefs will be only a function of the number of excess outcomes in one or the other direction. In particular, any sequence with balanced numbers leads to a correct uniform belief, which applies to, for example, both examples from the previous subsection: 000111 and 010101.

It's when both $\alpha > 0$ and $0 < \delta < 1$ that both effects are activated and more complex predictions obtain from the model. It's immediate that, under such parameters, $q_0(000111) > 0.5$, which coincides with the representativeness prediction. For the sequence $s^6 = 010101$, since the 0s are further back than the 1s, the prediction is that $q_0(010101) > 0.5$, which goes in the opposite direction to the representativeness prediction.

The two theories presented in this section, then, make different predictions for certain sequences, opening up the possibility of testing and differentiating between them using experimental data.

⁵This is the same form, with adjusted notation, seen in Benjamin et al. (2017, p. 15).

3 Experiment

The experimental design had the main goal of producing a rich enough dataset that could be used to evaluate the performance of each one of the theories from the previous section as well as to compare them to each other. Beyond this, it also had the goal of improving upon previous experiments, which have mainly focused on three paradigms:⁶

- i. Prediction Tasks: Participants are presented with a sequence of outcomes and asked to choose which outcome they think will happen next. This is the most straightforward test of the gambler’s fallacy, but the analyses of these tasks have focused on choices after a streak of the same outcome, ignoring more complex patterns of interaction. Furthermore, the streaks are often short, such as two, three or four repetitions.
- ii. Sequence Construction Tasks: Participants are instructed to generate sequences that they believe are random, in the sense of being most like what they would expect tosses of a fair coin to generate. These constructed sequences are then analyzed to check for frequencies of reversals and other statistical properties, and it’s also possible to consider which sequences people avoid constructing.
- iii. Randomness Judgment Tasks: Participants evaluate the randomness of given sequences, such as by being told that it may or may not have been generated by a fair coin and then being asked which case they think it is. Sequences can then be categorized according to such answers.

The main difficulty with tasks such as these is that there is not correct answer, since a fair coin is equally likely to produce any future outcome or any combination of outcomes. This complicates the issue of properly incentivizing responses, as well as judging whether such responses imply that subjects are truly wrong. The main way to fix this issue is to ask for *belief* responses. That is, how likely, in probability

⁶See Oskarsson et al. (2009, p. 264) for examples of each.

terms, subjects think each outcome is to come next. In this case there is one single objectively correct answer.

Having enough data richness to compare the two theories meant being capable of taking the analysis to the level of individual sequences. In order to achieve this, a large degree of randomization was employed in conjunction with a large number of tasks per subject. Subjects were presented with sequences of outcomes generated by a computerized version of a fair coin. These outcomes in these sequences were completely randomized, and their lengths varied from zero to eight outcomes, in increments of two, with the order of their appearance also being randomized. For each sequence, both choice and belief responses were collected.

3.1 Task

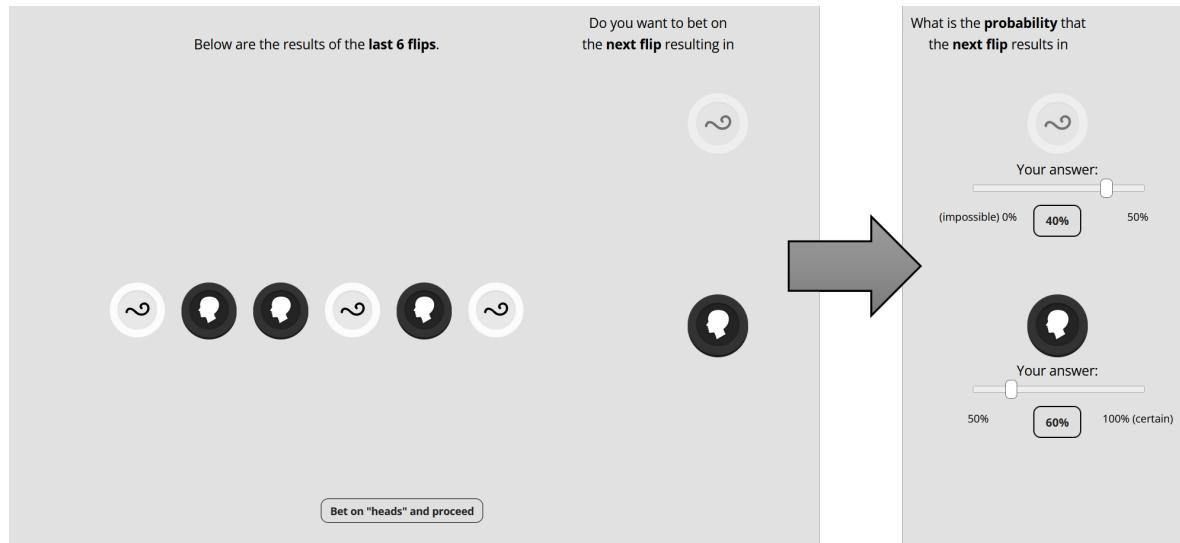


Figure 1: Typical task screen for a given round. Subjects are shown sequences of heads and tails and are asked for the next outcome and probabilities. The center-left of the screen shows the sequence of outcomes and remains the same throughout the round. The right side of the screen first asks for choice response, and then for probability responses. The button at the bottom of the screen is used to confirm both responses.

Figure 1 shows what a typical task screen looked like. In the center-left of the

screen, a sequence of heads and tails was presented to the subject.⁷ They then had to provide two responses by interacting with the right side of the screen and the confirmation button at the bottom. First, they had to click on which outcome, head or tail, they wanted to predict as being the next one. After that, sliders appeared beneath each outcome. By moving the sliders, they recorded their probabilistic beliefs, in integers from 0 to 100, regarding the likelihood of each outcome being the next one. When one slider was moved, the other one adjusted automatically. The sliders were programmed such that the probability for the chosen outcome had to be at least 50%. After the responses, a new outcome was generated to check if it matched the prediction. No feedback was provided.

3.2 Structure

After consenting to participate, subjects were presented with instruction screens explaining the nature of the task to them and how their answers would affect their bonus payment. Subjects were explained that the computer was producing the sequences they would be seeing in a way that exactly mirrored a real fair coin. The instructions first presented an explanation of the choice response, which was followed by 6 ‘warm up’ questions, in which subjects only had to choose the next outcome. Out of these 6 questions, one was chosen and if their prediction matched the actual result, \$0.60 was added to their bonus payment.⁸ Then the probability response was introduced and explained to them. At the end of the instructions, there was one practice question with both choice and probability responses just like a normal round. The bonus structure for choice and probabilities was explained to them and is detailed in the next subsection.

Subjects went through 50 rounds in which they encountered, in random order, 2 sequences of length 0⁹, 4 sequences of length 2, 12 sequences of length 4, 20 sequences

⁷A tail outcome was coded as 0, a head outcome was coded as 1. So, for example, the sequence appearing in figure 1 is 011010.

⁸This effectively amounted to an extra \$0.30 fixed payment in expectation regardless of their actual choices and was only meant to fix their understanding of the task before introducing belief responses.

⁹The question language was similar but slightly changed from the one seen in figure 1 for length 0.

of length 6 and 12 sequences of length 8. In each of these sequences, each outcome was randomized, i.e. it was equally likely to be a head or a tail. After these 50 sequences, subjects encountered 4 more sequences, one of each of the previous lengths and in random order. Each one of these sequences was composed of either all heads or all tails, each case being equally likely. For compensation purposes, these last 4 sequences were treated just like the previous 50, and this was explained to subjects.¹⁰

Before finishing the experiment, subjects were asked one non-incentivized multi-choice followup question asking for their reasoning in responding to the full run sequences. They were presented with 6 possible answers. See figure 2.

Follow-up Question

This question does not affect your approval or payment in any way.
But it will help us understand your answers better, so we ask your help in answering honestly.

Please choose the option that most closely approximates your reasoning when you were answering questions with sequences of flips like the one above:

- I do not trust the instructions about the coin/I believe there is deception in this study.
- When tossing a normal coin, streaks like this are more likely to continue.
- Previous flips like these are indicative of the bias of the coin, and the next flips should follow that.
- A normal coin should generate balanced outcomes, so future flips are more likely to break this streak than continue it.
- Previous flips do not affect future flips.
- I did not follow any specific reasoning.

Continue

Figure 2: Follow-up question at the end of the experiment asking subjects for their reasoning for their answers when encountering full runs of heads or tails.

3.3 Implementation

297 subjects were recruited through the Prolific website, an online platform specialized in maintaining a pool of subjects for scientific research. Using this platform allowed

Subjects were asked ‘Do you want to bet on the **first flip** resulting in’ and ‘What is the **probability** that the **first flip** results in’.

¹⁰Given that these last four sequences were not randomized in the same way as the previous 50, the language was adjusted to account for this. Subjects were asked what would their answers “have been” if they “had seen” these sequences.

for easy recruiting of a large number of participants from a broad population. Subjects were restricted to being located in the USA and being fluent in English.

Throughout the instructions, subjects were asked 4 simple comprehension questions about the functioning of the task and what the responses meant. In accordance with Prolific’s policies on comprehension questions, subjects who answered the same question incorrectly twice were not permitted to continue. This happened to 18 subjects.¹¹

The fixed payment for finishing the experiment was \$2.70. The bonus payment varied from \$0.00 to \$7.00. A potential bonus of \$0.60 came from the warm-up questions. Two questions out of the 54 were randomly selected for the remaining bonus. The choice response was worth a \$1.00 bonus. The probability response was worth a \$2.20 bonus, calculated according to the Binarized Scoring Rule (Hossain and Okui 2013). Following recent results from Danz et al. (2022), subjects were not directly given explicit mathematical formulas for this calculation, but were told that their bonus payment was maximized by truthfully reporting their best guesses for the correct probabilities. Mathematical formulas were provided at the end of the experiment.

The experiment was coded in JavaScript using custom plug-ins for the jsPsych library (de Leeuw 2015).

4 Results

In this section, I present the results of the analysis of the data from the experiment described in the previous section. I proceed from the aggregate to the granular level. I begin by showing that overall results from full runs of heads and tails look as one would expect from the existing literature. I then proceed to show, however, that there is significant heterogeneity among subjects, and they can be classified into four different groups: *rational*, *gambler’s fallacy*, *hot hands* and *both/mixed*. The analysis then proceeds by treating these groups separately, focusing on subject types

¹¹This is in addition to the 297 who successfully finished the experiment, resulting in a rate of 5.71%.

‘gf’ and ‘hh’, as these are the types for which there are significant results.¹² The first step proceeds at an aggregate level to verify the overall performance of each of the two models presented in section 2. Both perform well, with highly significant parameters. Once the analysis is taken to the level of individual sequences, I show that the Rabin and Vayanos model performs better than the entropy theory. Given this superior performance, I then verify whether increasing the number of parameters in a generalized version of the Rabin and Vayanos model increases performance. The general parameters are very close to the implied parameters from the simple model and offer no increase in explanatory power.

4.1 Basic Results

Figure 3 follows the basic structure of several figures appearing in this section. It shows both mean probability (left, on a 0-100 scale) and choice frequency (right, on a 0-1 scale) data for a group of sequences. The sequences displayed are those of full runs of heads or tails for length 2, 4, 6 and 8.¹³ The responses are for the outcome that continues the run, for example, frequency of choosing head as the next outcome when the sequence is a full run of heads. The responses are from all subjects.

Concentrating first on the choice data, it shows a significant gambler’s fallacy effect. Subjects are significantly more likely to choose the outcome that breaks the run than the one that continues it. Probability responses are less definitive, but there’s still a significant gambler’s fallacy effect for lengths 4 and 6: subjects believe the outcome that continues the run is less likely to happen than the one that breaks it. As the next subsection will show, however, looking at all subjects in this way obscures important differences between them. Nonetheless, this basic result shows that choice data, at the aggregate level, might overestimate the gambler’s fallacy effect. This further highlights the importance of the probability data collected in my experiment.

¹²Aggregate analysis for all subjects and the other subject types is provided in appendix B.

¹³Except when noted, in these figures, symmetric sequences (e.g. 0110 and 1001) are pooled together.

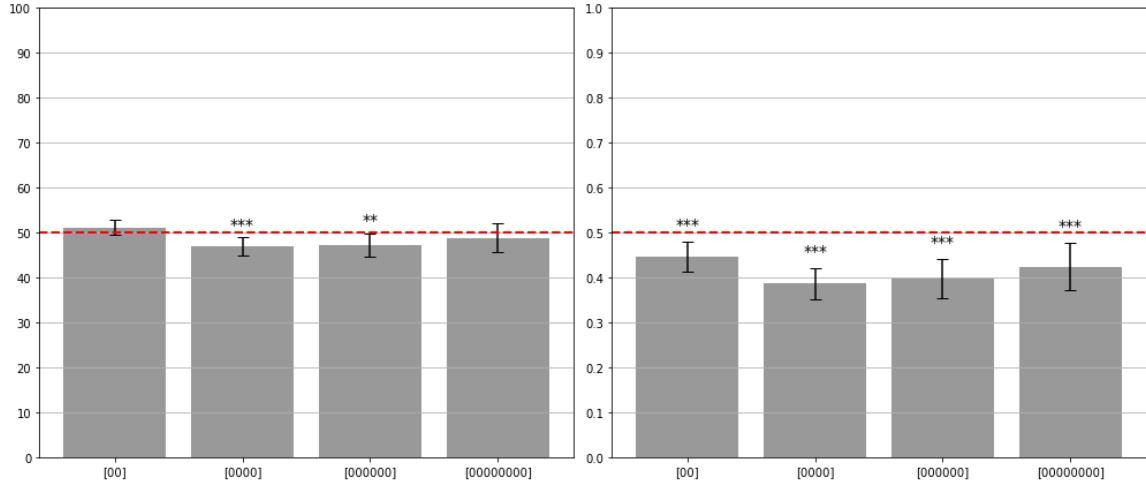


Figure 3: Mean probability (left) and choice frequency (right) responses for the outcome that continues the full run. All subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

4.2 Subject Heterogeneity

In the last four rounds of the experiment, sequences subjects encountered always took the form of full runs of either heads or tails, one of each of the lengths 2, 4, 6 and 8. Based on their probability responses for the outcome that continued the run in these rounds, subjects were classified into one of four types: '*rational (rat)*', '*gambler's fallacy (gf)*', '*hot hands (hh)*', '*both*'. Table 1 shows the behavior that led to being classified into each one.

Type	Classifying Behavior - Probability response for run-continuing outcome
rat	50% on all four rounds
gf	weakly below 50% on all four rounds, strictly below on at least one round
hh	weakly above 50% on all four rounds, strictly above on at least one round
both	strictly above 50% on at least one round and strictly below 50% on at least one round

Table 1: How subjects were classified into each of the four types depending on their probability responses for the outcome that continued the full runs in the final four rounds.

Based on these criteria, table 2 shows the distribution of types. The most common

type of subject was the ‘*rational*’ type, encompassing almost 30% of subjects. It was closely followed by the ‘*gambler’s fallacy*’ type. The least common type was the ‘*hot hands*’ type.

Type	Number	Proportion
rat	88	29.62%
gf	85	28.61%
both	66	22.22%
hh	58	19.52%

Table 2: Distribution of subject types.

As explained in section 3.2 and shown in figure 2, at the end of the experiment subjects were asked to choose between 6 options that best matched the reasoning for their responses to full run sequences. These 6 options were, in order, labeled as follows: (i) “I do not trust the instructions about the coin/I believe there is deception in this study.” as labeled ‘mistrust’; (ii) “When tossing a normal coin, streaks like this are more likely to continue.” was labeled ‘hh’; (iii) “Previous flips like these are indicative of the bias of the coin, and the next flips should follow that.” was labeled ‘inference’; (iv) “A normal coin should generate balanced outcomes, so future flips are more likely to break this streak than continue it.” was labeled ‘gf’; (v) “Previous flips do not affect future flips.” was labeled ‘rat’; (vi) “I did not follow any specific reasoning.” was labeled ‘*shrugs*’.

Figure 4 shows the distribution of these answer for each of the four types of subjects. The answers are mostly strongly in line with what would be expected, especially for ‘rat’ (figure 4a) and ‘gf’ (figure 4b) type subjects. About 85% and 80% of subjects of these types, respectively, gave the justification that matches their type. For ‘hh’ (figure 4c) type subjects, the three most common answers were those that could justify a belief that a full run of heads or tails is more likely to continue than to stop.

Figures 5, 6, 7 and 8 repeat figure 3 but considering individually each of the four types of subject: in order, ‘gf’, ‘hh’, ‘rat’ and ‘both’. These figures show that the

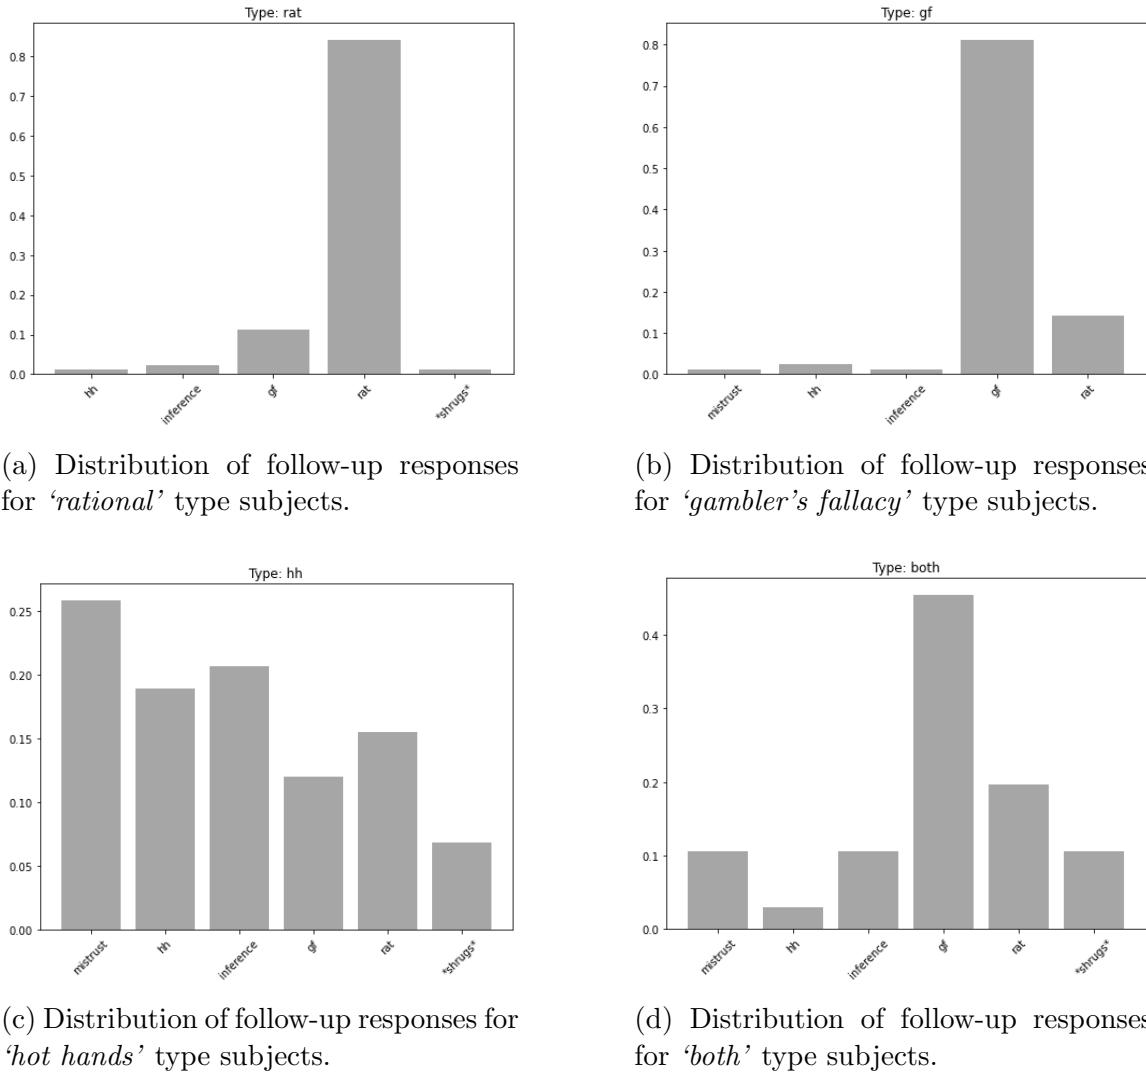


Figure 4: Distribution of follow-up responses for different subject types.

original aggregated data shown in the latter figure contained trends going in different directions. Types ‘gf’ (figure 5) and ‘hh’ (figure 6) show opposite tendencies: they expect the next outcome to be more likely to, respectively, reverse or continue the run, and this likelihood is monotone in the length of the run. Subjects of type ‘rat’ (figure 7) display an interesting pattern: their beliefs match the correct distribution, but their choices exhibit a gambler’s fallacy effect.¹⁴ Naturally, it’s not possible to

¹⁴Reminding the reader that both all-heads and all-tails sequences are included, so any systematic bias in the direction of one or the other outcome would be washed out by randomization.

affirmatively says that these subjects are ‘wrong’ in their choices for the reasons already discussed.¹⁵ Finally, for subjects of type ‘both’ (figure 8), it’s not really possible to offer any conclusive analysis. There’s a significant gambler’s fallacy effect for the sequences of length 4, but this effects disappears for longer lengths, and it’s not there for length 2 either. It’s possible to speculate that these subjects might switch their behavior depending on the length of the run, but the data as it stands cannot be used to conclusively justify such a conjecture.

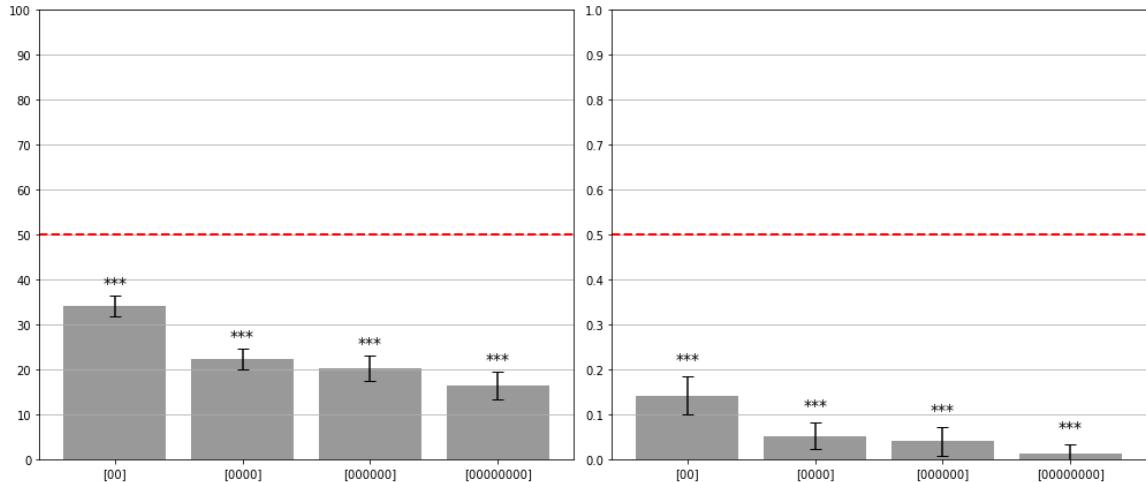


Figure 5: Mean probability (left) and choice frequency (right) responses for the outcome that continues the full run. ‘Gambler’s fallacy’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

4.3 Fitting the Rabin and Vayanos Model

Given the results from the previous subsection, the analysis now proceeds along lines that differentiate between different types of subjects, especially between ‘gf’ and ‘hh’ types. In this subsection, the Rabin and Vayanos (2010) model presented in subsection

¹⁵Given that the 50-50 response was always available, the restriction of choice and beliefs being consistent was non-binding for these subjects.

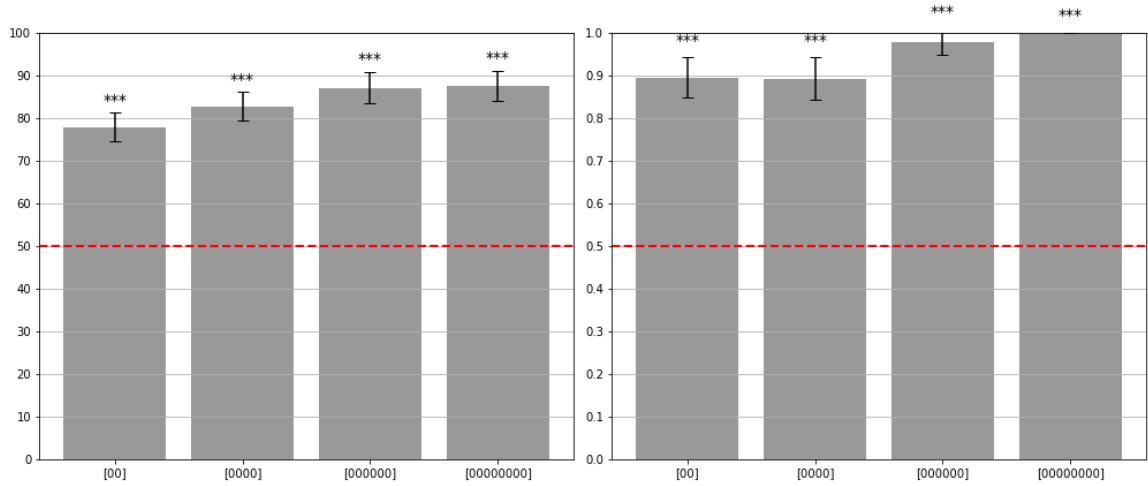


Figure 6: Mean probability (left) and choice frequency (right) responses for the outcome that continues the full run. ‘Hot hands’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

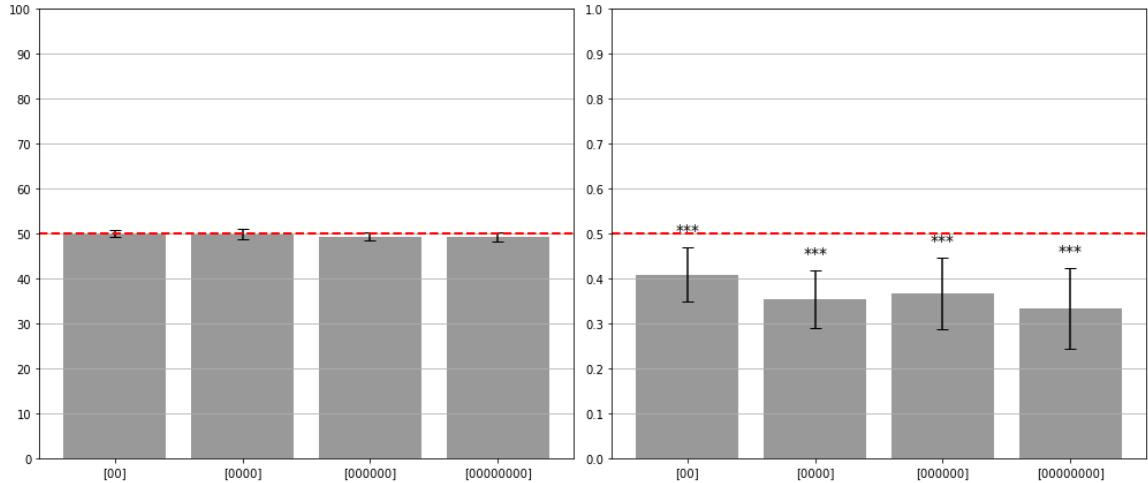


Figure 7: Mean probability (left) and choice frequency (right) responses for the outcome that continues the full run. ‘Rational’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

2.3 is fitted to the data for all sequences and responses from ‘gf’ type subjects, while controlling for sequence length and round.¹⁶ The fitted model is then given by

¹⁶As shown in subsection 4.6, for ‘hh’ type subjects, this is not the correct functional form, so an appropriate analysis for this type is delayed until then.

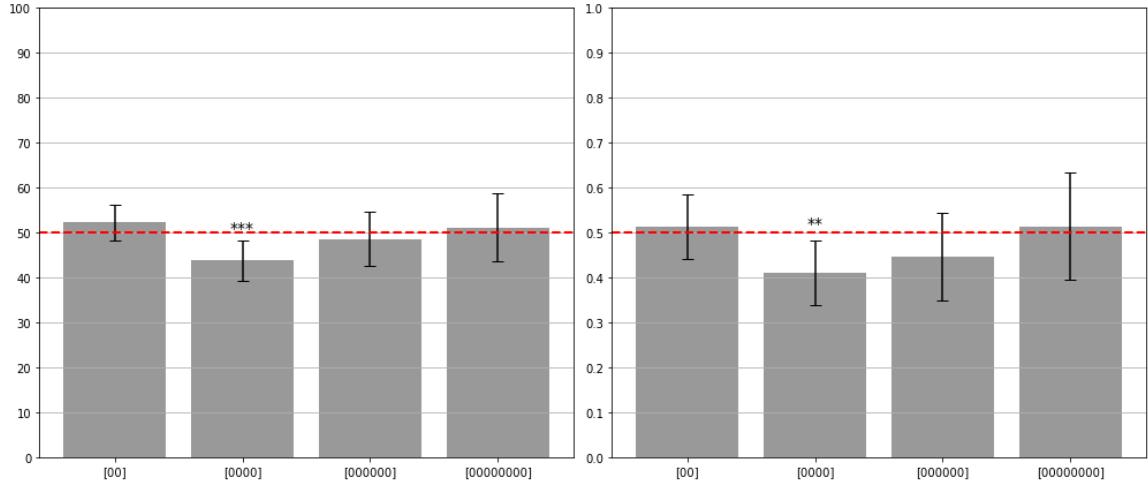


Figure 8: Mean probability (left) and choice frequency (right) responses for the outcome that continues the full run. ‘Both’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

$$q_{0j} = \sum_{i=1}^8 (\alpha \delta^i pos_{ij}) + \beta_{length} length_j + \beta_{round} round_j + \varepsilon_j \quad (1)$$

where q_0 is the reported probability of 0 being the next outcome, transformed to $[-1, 1]$, $length \in \{0, 2, 4, 6, 8\}$, $round \in \{1, 2, \dots, 53, 54\}$, and $pos_1, pos_2, \dots, pos_8$ are tertiary dummies encoding the value of the outcomes in each position according to

$$pos_i = \begin{cases} 1, & \text{if } i = 1, \\ 0, & \text{if } i = \emptyset, \\ -1, & \text{if } i = 0. \end{cases}$$

Table 3 reports the estimated parameters from a nonlinear least squares estimation. Both the reversal parameter $\alpha = 0.224$ and the recency parameter $\delta = 0.795$ are of a significant size, suggesting that both effects play an important role in subjects’ responses. This is only an analysis at an aggregate level, and a closer inspection in subsection 4.5 will indeed confirm this and provide additional details.

Parameter	Estimate
α	0.224324*** <i>(0.015481)</i>
δ	0.795059*** <i>(0.021439)</i>
β_{length}	-0.000587 <i>(0.001557)</i>
β_{round}	-0.000026 <i>(0.000281)</i>
<i>N</i>	4590

Table 3: Results for nonlinear least squares estimation of the Rabin and Vayanos (2010) model for ‘gf’ type subjects according to equation 1. Standard errors clustered at the subject level. * significant at 10% level, ** significant at 5% level, *** significant at 1% level.

4.4 Empirical Entropies

This subsection follows the theoretical results from subsection 2.2 to investigate the effect of empirical entropies on reported beliefs. In order to do so, a regression is run in order to obtain linear coefficients for the delta empirical entropies $\Delta_0 H_k$, $k = 0, 1, 2$, on reported beliefs of the next outcome being 0, transformed to $[-1, 1]$, while controlling for sequence length and round. All rounds are used in the regression. This is done separately for subjects of types ‘gf’ and ‘hh’. The regression equation is given by

$$q_{0j} = \alpha + \beta_0 \Delta_0 H_{0j} + \beta_1 \Delta_0 H_{1j} + \beta_2 \Delta_0 H_{2j} + \beta_3 length_j + \beta_4 round_j + \varepsilon_j$$

Remember that $\Delta_0 H_k(s^n) = H_k(s^n0) - H_k(s^n1)$ and that, for the gambler’s fallacy, it is expected that the outcome that makes the sequence more complex are more likely. Since the dependent variable is the probability of 0, for ‘gf’ type subjects, the expected signs of the coefficients are $\beta_0 > 0$, $\beta_1 > 0$ and $\beta_2 > 0$. If ‘hh’ type subjects are the mirror of ‘gf’ subjects, one would expect that $\beta_0 < 0$, $\beta_1 < 0$ and $\beta_2 < 0$.

Table 4 reports the results of these regressions. For ‘gf’ type subjects, the effects of empirical entropies are all highly significant, and their effect goes in the expected direction. The larger the increase in complexity that outcome 0 makes when com-

pared to outcome 1, the more likely ‘gf’ subjects think outcome 0 is to come next. Another thing to observe from the results is that $\beta_0 > \beta_1 > \beta_2$, meaning that lower order empirical entropies have a larger effect on beliefs.¹⁷ This is to be expected if one considers that higher order entropies capture more complicated patterns, which subjects might find harder to notice or reason about.

For ‘hh’ type subjects, the coefficients also have the expected sign. It’s still the case that the coefficient for order zero is the largest one, but the coefficient for order one is not as significant and is the smallest one. This continues a broader trend of ‘hh’ type subjects not being perfect mirrors of ‘gf’ type subjects.

	q_0	
	(gf)	(hh)
<i>const</i>	-0.0259 (0.016)	-0.0180 (0.031)
$\Delta_0 H_0$	0.5856*** (0.037)	-0.7366*** (0.059)
$\Delta_0 H_1$	0.2621*** (0.036)	-0.1186* (0.062)
$\Delta_0 H_2$	0.1254*** (0.044)	-0.2536*** (0.076)
<i>length</i>	0.0026 (0.003)	-0.0011 (0.004)
<i>rounds</i>	0.0002 (0.000)	0.0009 (0.001)
<i>R</i> ²	0.308	0.268
N	4590	3132

Table 4: Regression results for ‘gf’ and ‘hh’ type subjects. q_0 , the dependent variable, is the reported probability of 0 as the next outcome. OLS regression with standard errors clustered at the subject level. * significant at 10% level, ** significant at 5% level, *** significant at 1% level.

¹⁷The null hypothesis that $\beta_0 \leq \beta_1$ is rejected at the 1% level, and the null hypothesis that $\beta_1 \leq \beta_2$ is rejected at the 5% level.

4.5 Closer Look: Individual Sequences

The previous two subsections have shown that at an aggregate level, both theories presented in section 2 perform well when explaining the gambler’s fallacy effects found in subjects’ behavior. The Rabin and Vayanos (2010) model results in significant coefficients that confirm the reservation and recency effects. While empirical entropies have been shown to have a highly significant effect on beliefs, with outcomes that increase sequence complexity being seen as more likely. This subsection, therefore, moves the analysis from an aggregate level down to the level of individual sequences in order to take a closer look at subject behavior and allow the differentiation between the two theories.

Appendix A provides the full look at all sequences of length 2, 4 and 6. This subsection focuses on the cases that provide the sharpest differentiation between the two theories. In particular, it presents the case of the *alternating sequences* (0101... and 1010...), for which subjects predict a continuation of the pattern rather than its break. This violates the predictions made by the entropy model, while still closely following the predictions of the Rabin and Vayanos (2010) model. It also presents the case of the sequences for which both outcomes affect complexity equally, and so no prediction comes from the entropy model. Nonetheless, subjects still, for the most part, deviate from correct beliefs. The most important effect driving these results is the recency effect created by the discount factor δ , and so this effect is isolated for the cases of sequences with just one different outcome, but in different positions.

The figures in this subsection are similar to the ones shown in subsection 4.1. Probability responses are presented on the left and choice responses on the right. Except when noted, these responses are for the outcome that results in lower entropy. The red dots shown on the left panels mark the point-prediction of the Rabin and Vayanos model for the sequence when using the parameters estimated in subsection 4.3 ($(\alpha, \delta) = (0.224, 0.795)$). All results are for ‘gf’ type subjects only.

4.5.1 Length 4: $\Delta H_0 = 0$, $\Delta H_1 > 0$, $\Delta H_2 = 0$

Figure 9 shows responses for low entropy outcome for length 4 sequences that have $\Delta H_0 = 0$, $\Delta H_1 > 0$, $\Delta H_2 = 0$. It is among these sequences that some of the most significant violations of the entropy theory happen. For sequences 0101, 1010, 0110 and 1001, subjects believe that the low entropy outcome is more likely to come next. While these responses violate the entropy prediction, they are in line with the prediction made by the Rabin and Vayanos model. For example, for the alternating sequences, 0101 and 1010, because of the recency effect, one of the outcomes is further to the front in all instances and so gets overweighted relative to the other one, leading subjects to believe it less likely to come next.

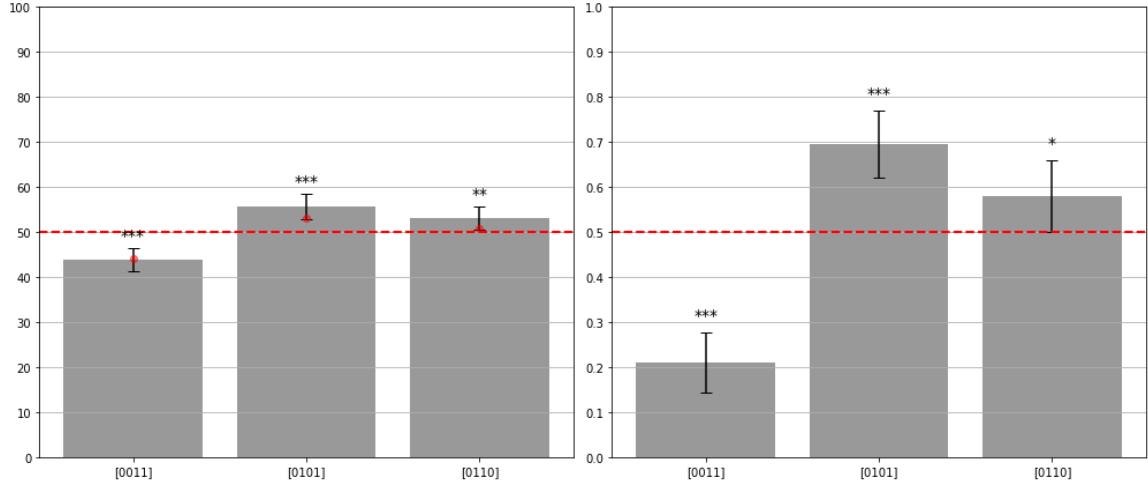


Figure 9: Mean probability (left) and choice frequency (right) responses for lower entropy outcome for length 4 sequences with $\Delta H_0 = 0$, $\Delta H_1 > 0$, $\Delta H_2 = 0$. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

4.5.2 Length 6: $\Delta H_0 = 0$, $\Delta H_1 > 0$, $\Delta H_2 > 0$

Figure 10 shows responses for low entropy outcome for length 6 sequences that have $\Delta H_0 = 0$, $\Delta H_1 > 0$, $\Delta H_2 > 0$. Although at a lower significance level, the alternating sequence here further confirms the same effect as seen in the previous figure. Subjects

responses go in the opposite direction to that predicted by the entropy model, while being in line with the prediction of the Rabin and Vayanos model.

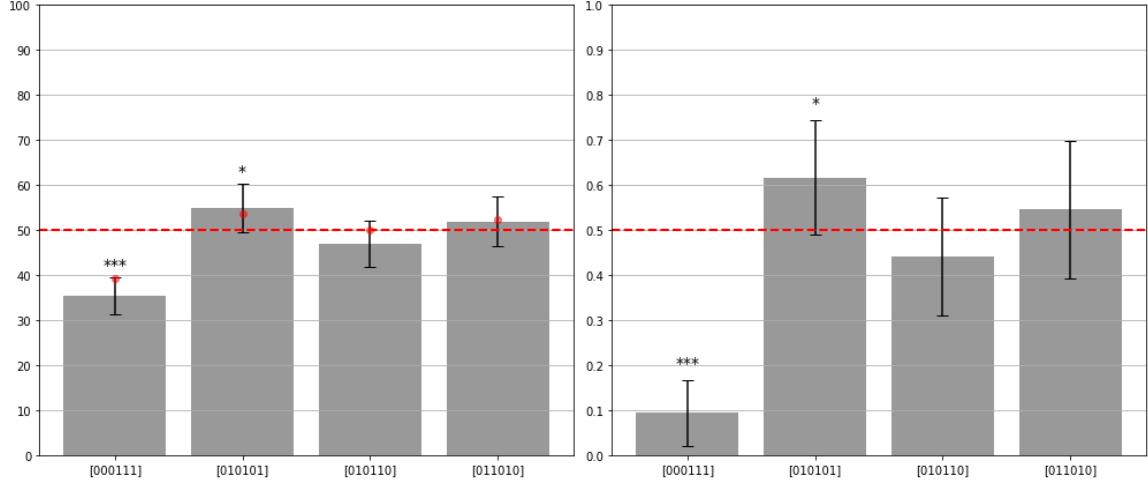


Figure 10: Mean probability (left) and choice frequency (right) responses for lower entropy outcome for length 6 sequences with $\Delta H_0 = 0$, $\Delta H_1 > 0$, $\Delta H_2 > 0$. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

4.5.3 Length 6: $\Delta H_0 = \Delta H_1 = \Delta H_2 = 0$

Figure 11 shows responses for outcome 0 for length 6 sequences that have $\Delta H_0 = \Delta H_1 = \Delta H_2 = 0$. That is, these are sequences for which there is no ‘low entropy’ outcome, as both outcomes 0 and 1 result in a sequence with same empirical entropy of orders 0, 1 and 2. These can be thought of sequences that are already highly complex, and there’s little pattern to be matched or reinforced by either outcome. Nevertheless, it’s shown that for most of these sequences, subjects still deviate significantly from the correct belief. Their deviation closely follows the prediction of Rabin and Vayanos. It’s worth noting that this happens with sequences that have the same outcome repeated twice at the end of the sequence, further reinforcing the importance of the recency effect.

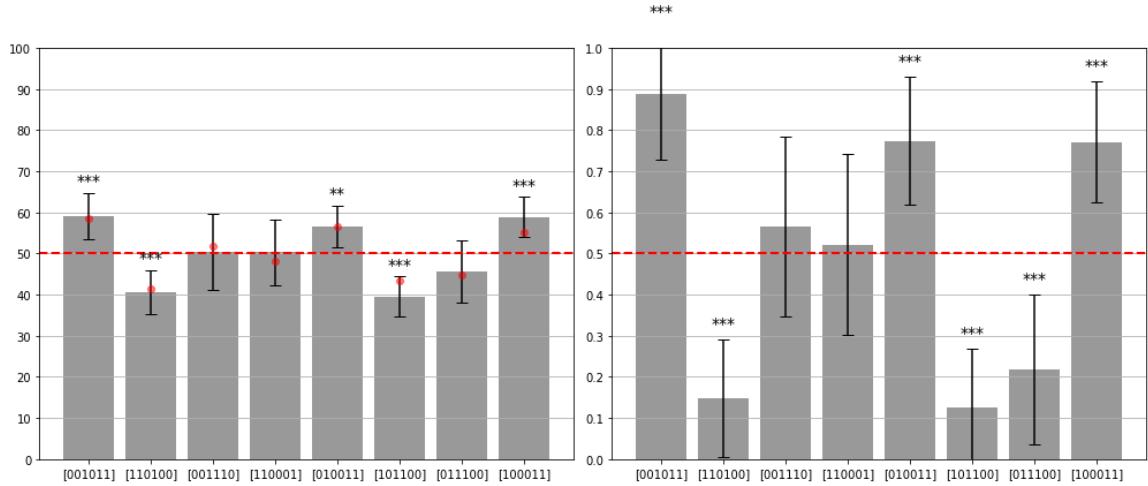


Figure 11: Mean probability (left) and choice frequency (right) responses for outcome 0 for length 6 sequences with $\Delta H_0 = 0$, $\Delta H_1 = 0$, $\Delta H_2 = 0$. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

4.5.4 Lengths 4 and 6: recency effects

Figures 12 and 13 illustrate in more detail the recency effect for sequences of length 4 and 6, respectively. They show in order sequences that have one (or no) outcome different than the rest, but differing by the position in which it appears. These figures show a monotonicity that is created by the recency effect. The closer the different outcome is to the front of the sequence, the larger it’s impact in countering the effect of the rest of the outcomes on the belief. For example, subjects believe 0 is less likely to be next when facing the sequence 000001 than the sequence 100000. This cannot be explained by empirical entropy, but is in line with the Rabin and Vayanos model.

4.6 Generalized Rabin and Vayanos

The previous subsection has shown that the Rabin and Vayanos (2010) model has a superior performance in matching the actual behavior of ‘gf’ type subjects. This subsection explores whether it is possible to increase the performance of the model in a significant way by increasing the number of parameters, creating a generalized

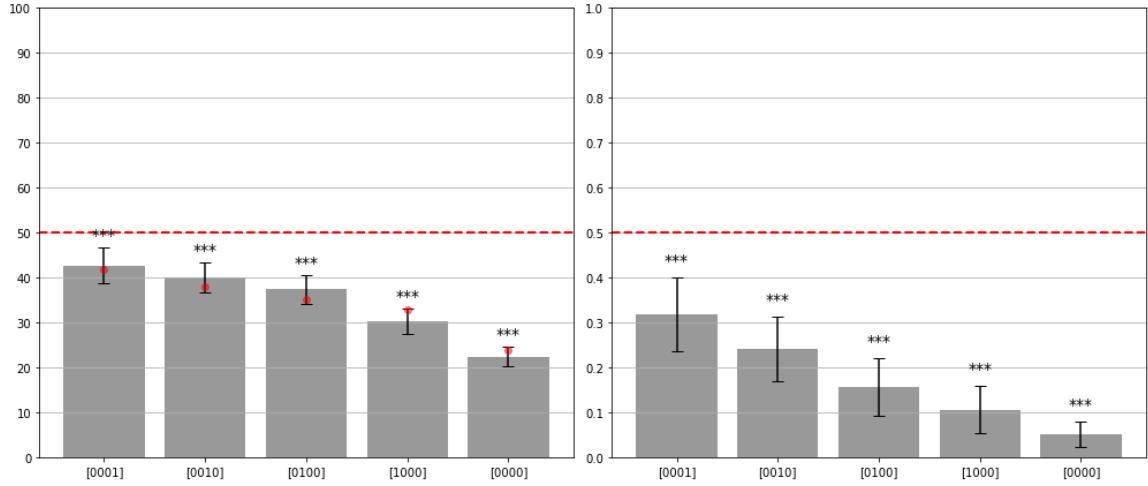


Figure 12: Mean probability (left) and choice frequency (right) responses for lower entropy outcome, showing recency effect for length 4 sequences. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

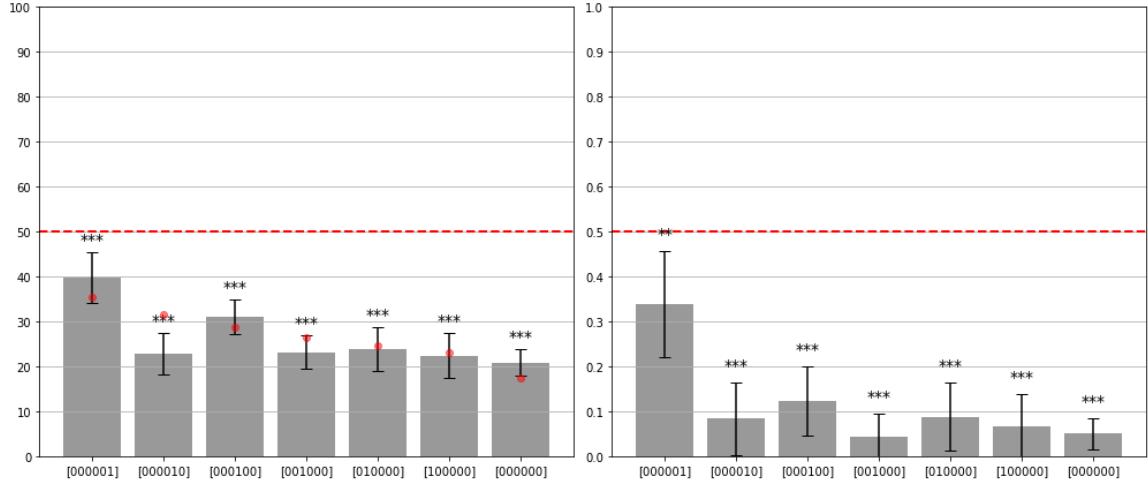


Figure 13: Mean probability (left) and choice frequency (right) responses for lower entropy outcome, showing recency effect for length 6 sequences. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

version of their model.

The generalized version of the model takes the form

$$q_{0j} = \sum_{i=1}^8 (\beta_i pos_{ij}) + \beta_{length} length_j + \beta_{round} round_j + \varepsilon_j. \quad (2)$$

Before, the model had only two parameters, α and δ , and it was the case that they implied $\beta_i = \alpha\delta^i$. These implied parameters can be compared to the ones that are estimated separately in the generalized model. Table 5 reports the results of this estimation for subject types ‘gf’ and ‘hh’. The position parameters are all highly significant and with the expected signs.

Figure 14 shows the comparison of the estimated values of the general parameters and the parameters implied by the previously estimated α and δ for ‘gf’ type subjects. The conclusion is that the generalization makes little difference. The free estimation of the 8 parameters result in numbers that are very close to the numbers that were already implied by α and δ . This conclusion is further reinforced by a likelihood ratio test between the two models, which returns a p-value of 0.169, failing to reject even at the 10% significance level the null hypothesis that the 8 parameter model provides no additional explanatory power compared to the 2 parameter model.

Figure 15 plots the estimated parameters for ‘gf’ and ‘hh’ subjects. It highlights the difference in behavior for ‘hh’ subjects. While for ‘gf’ subjects the parameters show the recency bias implied by the discount factor, with more recent results being weighted more heavily, for ‘hh’ subjects, it’s actually the outcomes in the middle of the sequence that are more heavily weighted.

5 Continuing the Gambler’s Fallacy

Imagine yourself, once again and against all odds, observing someone toss a fair coin ten times in front of you, resulting in the same sequence: head, head, head, head, head, head, head, head, head, head. But now you are then asked which outcome is more likely if the coin is tossed twice again: a head and another head, a head and then a tail, a tail and then a head, or a tail and then another tail? Even in this

	q_0	
	(gf)	(hh)
<i>const</i>	-0.0261 (0.016)	-0.0035 (0.029)
<i>pos</i> ₁	0.1849*** (0.011)	-0.0455** (0.021)
<i>pos</i> ₂	0.1344*** (0.010)	-0.1023*** (0.015)
<i>pos</i> ₃	0.1099*** (0.008)	-0.1377*** (0.012)
<i>pos</i> ₄	0.0800*** (0.010)	-0.1514*** (0.014)
<i>pos</i> ₅	0.0796*** (0.009)	-0.1165*** (0.014)
<i>pos</i> ₆	0.0646*** (0.008)	-0.1298*** (0.011)
<i>pos</i> ₇	0.0477*** (0.014)	-0.1035*** (0.022)
<i>pos</i> ₈	0.0425*** (0.012)	-0.0528*** (0.017)
<i>rounds</i>	0.0003 (0.000)	0.0010 (0.001)
<i>length</i>	0.0024 (0.003)	-0.0032 (0.004)
<i>R</i> ²	0.425	0.276
N	4590	3132

Table 5: Regression results for ‘gf’ and ‘hh’ type subjects of the generalized Rabin and Vayanos model. q_0 , the dependent variable, is the reported probability of 0 as the next outcome. OLS regression with standard errors clustered at the subject level. * significant at 10% level, ** significant at 5% level, *** significant at 1% level.

standard case of a sequence made up entirely of a single outcome, all heads or all tails, the answer from an agent suffering from the gambler’s fallacy is not clear. The gambler’s fallacy might lead to an expectation of, for example, a head after four tails, but if the next two outcomes will be two heads or a head and a tail cannot be so intuitively determined. On the one hand, the streak has already been broken, but on the other, there are still four tails in a row and a single head after that might not be considered enough to balance them out. And yet, two heads, one after the other, to better provide this balance, might be seen as its own issue.

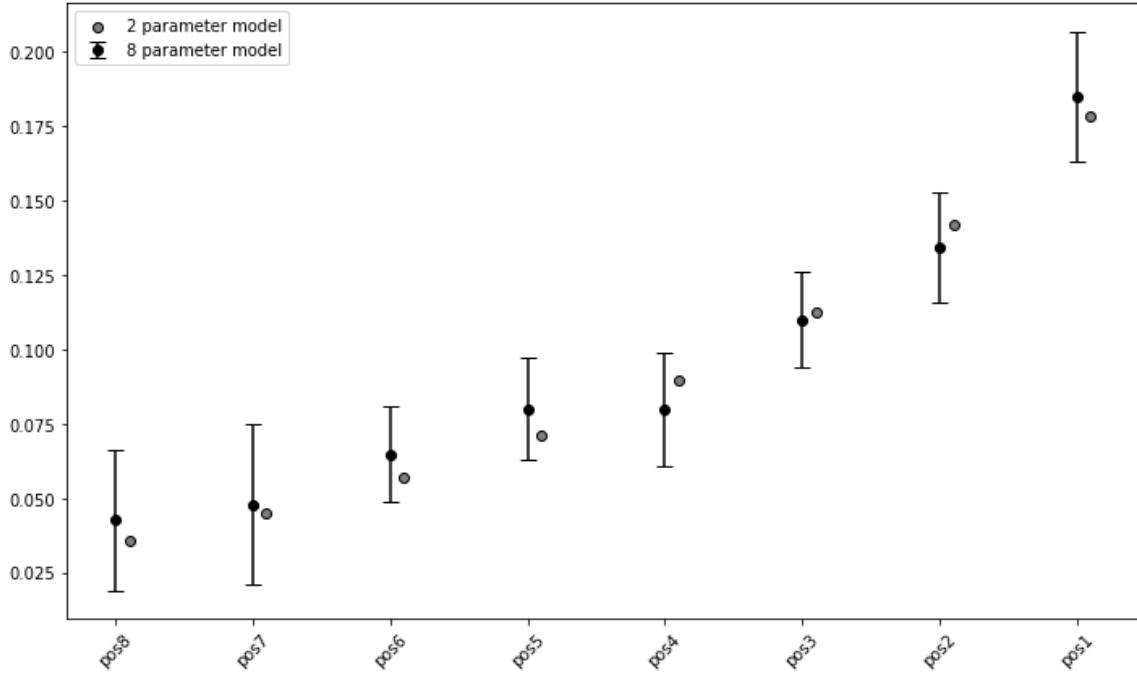


Figure 14: Comparison of the estimated 8 parameters in the generalized Rabin and Vayanos model and the implied 2 parameters original version. Error bars are 95% confidence intervals. ‘gf’ type subjects.

Given such ambiguities, this section expands the paradigm developed in previous sections to the case of predicting the next outcomes of the random process.

5.1 Theory

5.1.1 Notation

Similarly to the case of predicting the single next outcome, an agent faces finite binary strings of varying length $n \in \mathbb{N}$. Their notation is given by $s^n \in 2^\omega$. Particular interest is given to such sequences that have already been shown to be relevant in the previous case, for example $s^4 = 1111$ or $s^6 = 010101$. These strings are produced by independent draws from a distribution $\rho \in \Delta(\{0, 1\})$, and the agent is informed of this process. In the case of a fair coin, ρ is the uniform Bernoulli distribution. After seeing a sequence s^n , the agents reports a belief p over the next two outcomes. That is, the agent’s responses follow a belief mapping $s^n \mapsto p \in \Delta(\{00, 01, 10, 11\})$. The individual

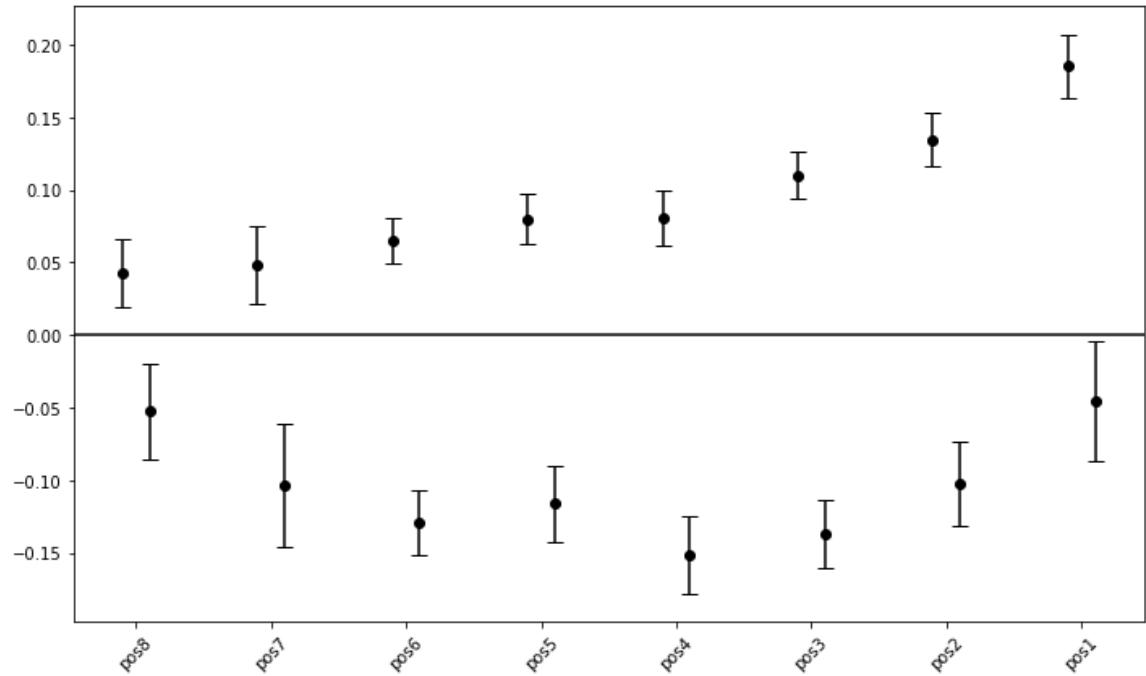


Figure 15: Comparison of estimated parameters of the generalized Rabin and Vayanos model for ‘gf’ type subjects (top) and ‘hh’ type subjects (bottom). Error bars are 95% confidence intervals.

beliefs over specific outcomes are denoted via subscripts, with $p_{00}, p_{01}, p_{10}, p_{11}$.

Notice that, given a belief over the next two outcomes for a sequence, $p(s^n)$, it is implied a belief of the next outcome being 0 or 1. This is denoted as $p_0(s^n) = p_{00}(s^n) + p_{01}(s^n)$ and $p_1(s^n) = p_{10}(s^n) + p_{11}(s^n)$. This *reduced belief* over the next outcome only can be studied much in the same way as has been done in the previous sections. For example, a classic gambler’s fallacy case is given by $p_0(0000) < 0.5$. Other possible reductions give beliefs over the second outcome only ($p_{00}(s^n) + p_{10}(s^n), p_{01}(s^n) + p_{11}(s^n)$) or over repeating the same outcome or switching ($p_{00}(s^n) + p_{11}(s^n), p_{01}(s^n) + p_{10}(s^n)$).

Finally, a notation of the form $s^n\nu$, with $\nu \in \{0, 1, 00, 01, 10, 11\}$, will define the concatenation of the original sequence s^n with ν at the end. For example, if $s^n = 1111$ and $\nu = 0$, then $s^n\nu = s^n0 = 11110$, and if $s^n = 0101$ and $\nu = 00$, then $s^n\nu = s^n00 = 010100$.

5.1.2 Extending Rabin and Vayanos

Recall the recency-weighted reversal model presented by Rabin and Vayanos (2010), which was shown to be very effective at describing belief behavior among subjects who were facing the problem of predicting the single next outcome. To extend the model to the two outcomes case, consider the probability of the first next outcome and the probability of the second next outcome conditional on the first next outcome. Hence, define:

$$p_{00}(s^n) = p_0(s^n)p_0(s^n0)$$

$$p_{01}(s^n) = p_0(s^n)p_1(s^n0)$$

$$p_{10}(s^n) = p_1(s^n)p_0(s^n1)$$

$$p_{11}(s^n) = p_1(s^n)p_1(s^n1)$$

That is, given a sequence s^n , to obtain the belief for outcomes ν , say $\nu = 01$, take the belief for the single outcome case for the next outcome in ν , 0, given the sequence s^n , and then multiply it by the belief for the single outcome case for the second outcome in ν , 1, given the sequence s^n and conditional on the next outcome being the first outcome in ν .

In general, both for the one next outcome and the two next outcomes cases, whether the resulting belief will deviate from the uniform distribution in one or another direction will depend on the exact values of the parameters. However, for the one outcome case, there are some important classes of sequences for which the prediction is unambiguous. One such class is that of the streaks of a single outcome, such as $s^4 = 0000$ and $s^n = 111111$. For such sequences, the one next outcome model delivers the standard gambler's fallacy effect of a higher probability for the outcome that breaks the streak, and this probability is monotone in the length of

the sequence, e.g., $p_0(00000) < p_0(0000) < 0.5$. Another class is that of alternating sequences of even length, such as $s^4 = 0101$ or $s^6 = 101010$, for which the one next model predicts that the alternating outcome is more likely to come next than the one that repeats the last outcome, and this is also monotone in the length of sequence. That is, $p_0(010101) > p_0(0101) > 0.5$.¹⁸

The case for the next two outcomes is not so immediate. As discussed intuitively in the introduction, for the streak sequences, whether one assigns higher likelihood to two of the outcome that breaks the streak or whether one expects it to be more likely first to break the streak but then switch again is not clear. Likewise, although it's reasonable to assume that in the alternating case the prediction is to reinforce the pattern continuation, this requires verification.

To make things easier to see, rewrite the expressions for the beliefs as follows:

$$p_{00}(s^n) = \frac{1}{4} \left[q_0(s^n)(1 + q_0(s^n0)) + q_0(s^n0) \right] + \frac{1}{4},$$

$$\begin{aligned} p_{01}(s^n) &= \frac{1}{4} \left[q_0(s^n)(1 - q_0(s^n0)) - q_0(s^n0) \right] + \frac{1}{4} \\ &= \frac{1}{4} \left[q_0(s^n)(1 + q_1(s^n0)) + q_1(s^n0) \right] + \frac{1}{4}, \end{aligned}$$

$$\begin{aligned} p_{10}(s^n) &= \frac{1}{4} \left[-q_0(s^n)(1 + q_0(s^n1)) + q_0(s^n1) \right] + \frac{1}{4} \\ &= \frac{1}{4} \left[q_1(s^n)(1 + q_0(s^n1)) + q_0(s^n1) \right] + \frac{1}{4}, \end{aligned}$$

¹⁸The monotonicity here being in terms of adding more alternating pairs, e.g., $s^2 = 01$, $s^4 = 0101$, $s^6 = 010101$ and so on.

$$\begin{aligned}
p_{11}(s^n) &= \frac{1}{4} \left[-q_0(s^n)(1 - q_0(s^n 1)) - q_0(s^n 1) \right] + \frac{1}{4} \\
&= \frac{1}{4} \left[q_1(s^n)(1 + q_1(s^n 1)) + q_1(s^n 1) \right] + \frac{1}{4}.
\end{aligned}$$

We can then see that whether the beliefs will deviate from the uniform case in one or another direction will depend on the sign of the expressions inside the brackets, and the ordering between them will depend on the relative values of these expressions. It also highlights the role of $q_0(s^n 0)$ and $q_0(s^n 1)$, which can be noticed to be such that $q_0(s^n 0) = \delta q_0(s^n) - \alpha\delta$ and $q_0(s^n 1) = \delta q_0(s^n) + \alpha\delta$. Indeed, using these expressions for $p_{00}(s^n)$, $p_{01}(s^n)$, $p_{10}(s^n)$, $p_{11}(s^n)$, $q_0(s^n 0)$ and $q_0(s^n 1)$ allows one to make the following observations.¹⁹

Observations 1. *In the extended Rabin and Vayanos model, the following hold for any sequence $s^n \in 2^\omega$:*

- i. $q_0(s^n 0) > 0 \Rightarrow p_{00}(s^n) > p_{01}(s^n)$ and $q_0(s^n 0) < 0 \Rightarrow p_{00}(s^n) < p_{01}(s^n)$,
- ii. $q_0(s^n 1) > 0 \Rightarrow p_{10}(s^n) > p_{11}(s^n)$ and $q_0(s^n 1) < 0 \Rightarrow p_{10}(s^n) < p_{11}(s^n)$,
- iii. $q_0(s^n) > 0 \Rightarrow p_{00}(s^n) > p_{11}(s^n)$ and $q_0(s^n) < 0 \Rightarrow p_{00}(s^n) < p_{11}(s^n)$,
- iv. $q_0(s^n) > 0 \Rightarrow p_{01}(s^n) > p_{10}(s^n)$ and $q_0(s^n) < 0 \Rightarrow p_{10}(s^n) < p_{01}(s^n)$
- v. $q_0(s^n) > \alpha\delta \Rightarrow p_{00}(s^n) > p_{10}(s^n)$
- vi. $q_0(s^n) > 0 \Rightarrow p_{01}(s^n) > p_{11}(s^n)$ and $(q_0(s^n) < 0 \text{ and } q_0(s^n) < -\alpha\delta) \Rightarrow p_{11}(s^n) > p_{01}(s^n)$
- vii. $q_0(s^n 0) \ R \ 0 \Leftrightarrow \sum_{t=0}^{\infty} \delta^{t+1} s_{n-t}^n \ R \ 1, \ R \in \{=, >, <\}$

¹⁹Algebraic manipulations using these expressions will show, for any $R \in \{=, >, <\}$, the equivalences: $p_{01}(s^n) \ R \ p_{00}(s^n) \Leftrightarrow q_0(s^n 0) \ R \ -q_0(s^n)q_0(s^n 0)$, $p_{10}(s^n) \ R \ p_{11}(s^n) \Leftrightarrow q_0(s^n 1) \ R \ q_0(s^n)q_0(s^n 1)$, $p_{00}(s^n) \ R \ p_{11}(s^n) \Leftrightarrow q_0(s^n)^2 \ R \ q_0(s^n)$, $p_{01}(s^n) \ R \ p_{10}(s^n) \Leftrightarrow q_0(s^n)(1 + \alpha\delta) \ R \ q_0(s^n)\delta$, $p_{00}(s^n) \ R \ p_{10}(s^n) \Leftrightarrow q_0(s^n)(1 + q_0(s^n 0)) \ R \ (1 - q_0(s^n))\alpha\delta$, $p_{11}(s^n) \ R \ p_{01}(s^n) \Leftrightarrow -\alpha\delta \ R \ q_0(s^n)(1 + \alpha\delta)$. The sufficient conditions in observations (i)-(vi) then follow either by implication or disjunctive syllogism. (vii) and (viii) follow from $q_0(s^n 0) = \delta q_0(s^n) - \alpha\delta$ and $q_0(s^n 1) = \delta q_0(s^n) + \alpha\delta$.

$$viii. q_0(s^n 1) R 0 \Leftrightarrow \sum_{t=0}^{\infty} \delta^{t+1} s_{n-t}^n R -1, R \in \{=, >, <\}$$

The observations above are useful in establishing necessary implications of the model regardless of specific parameter values. They link behavior of beliefs about the next outcome, which comes from the original model, to beliefs about the next two outcomes generated by the extended model. In particular, they can be used to show necessary predictions that the extended model makes about the important classes of sequences: the streak sequences and the alternating sequences.

Streak Sequences

Consider sequences s^n with length $n \geq 2$ given entirely by just 0s or just 1s, e.g., $s^2 = 11$ or $s^8 = 00000000$. Due to symmetry, let us focus on those with just 0s. The crucial question for these sequences is whether $p_{11}(s^n) > p_{10}(s^n)$ or vice-versa.

Before considering that, though, notice the following. Since $q_0(s^n) < 0$, by observations 1.iii and 1.iv, it's the case that $p_{11}(s^n) > p_{00}(s^n)$ and $p_{01}(s^n) > p_{10}(s^n)$. It's also the case that $q_0(s^n) < -\alpha\delta$, so by observation 1.vi, $p_{11}(s^n) > p_{01}(s^n)$. Finally, since $\sum_{t=0}^{\infty} \delta^{t+1} s_{n-t}^n < 0$, observation 1.vii and observation 1.i imply that $p_{01}(s^n) > p_{00}(s^n)$. Therefore, one of two orderings of beliefs must hold, depending on whether $p_{11}(s^n) > p_{10}(s^n)$ or $p_{11}(s^n) < p_{10}(s^n)$:

$$p_{11}(s^n) > p_{10}(s^n) > p_{01}(s^n) > p_{00}(s^n) \text{ or } p_{10}(s^n) > p_{11}(s^n) > p_{01}(s^n) > p_{00}(s^n)$$

From observation 1.ii, which of the two orderings must hold will depend on the sign of $q_0(s^n 1)$, which in turn by observation 1.viii hinges on $\sum_{t=0}^{\infty} \delta^{t+1} s_{n-t}^n < -1$ or not. And then, finally, notice that $\sum_{t=0}^{\infty} \delta^{t+1} s_{n-t}^n$ are the partial sums of the negative of the geometric series with ratio δ . Therefore, $\delta \leq 0.5$ immediately implies that it will always be $p_{10}(s^n) > p_{11}(s^n)$. That is, the recency effect is so strong that no matter how long the streak is, there's a greater concern with not having two 1s in a row at the end of the sequence.

For $\delta > 0.5$, on the other hand, there's an interaction between the value of δ and the length of the streak n . For a fixed length n , there's a $\tilde{\delta}(n)$ such that $p_{11}(s^n) > p_{10}(s^n)$ if $\delta \geq \tilde{\delta}(n)$. And for a given δ , the partial sum of the n terms of the series

will be greater than 1 if, and only if

$$\frac{\delta(1 - \delta^n)}{1 - \delta} > 1 \Leftrightarrow \delta^n < \frac{2\delta - 1}{\delta} \Leftrightarrow n > \frac{\ln(2\delta - 1)}{\ln(\delta)} - 1.$$

Thus, for a given $\delta > 0.5$, there's $N(\delta)$ given by the smallest positive integer such that $N(\delta) > \ln(2\delta - 1)/\ln(\delta) - 1$. For sequences shorter than $N(\delta)$, $p_{10}(s^n) > p_{11}(s^n) > p_{01}(s^n) > p_{00}(s^n)$, while for those whose length exceeds $N(\delta)$, it switches to the other ordering.

The one belief behavior which is prohibited by these results is having $p_{11}(s^n) > p_{10}(s^n)$ for shorter sequences and then switching to $p_{10}(s^n) > p_{11}(s^n)$ as streak length increases. Otherwise, it might be that $p_{11}(s^n) > p_{10}(s^n)$ for all streaks, $p_{10}(s^n) > p_{11}(s^n)$ for all streaks, or there's a switch from $p_{10}(s^n) > p_{11}(s^n)$ to $p_{11}(s^n) > p_{10}(s^n)$ as streak length increases.

For empirical purposes, however, it is important to consider how likely the switch from $p_{10}(s^n) > p_{11}(s^n)$ to $p_{11}(s^n) > p_{10}(s^n)$ actually is to happen. This depends on whether $N(\delta)$ is high enough, which in turn depends on δ not being too large. Now recall that in the results from the original experiment, for example, the estimated δ was $\delta \approx 0.79$, and for such we have $N(\delta) = 2$, hence no switch is expected. In order for a switch to be consistent with the model, a smaller δ would be needed.²⁰

Alternating Sequences

Consider now the sequences s^n of even length n which are alternating, e.g., $s^2 = 01$, $s^6 = 101010$. Again due to symmetry, focus on those starting with a 0 and ending in 1.

Because $q_0(s^n) > 0$, $p_{00}(s^n) > p_{11}(s^n)$, $p_{01}(s^n) > p_{10}(s^n)$, $p_{01}(s^n) > p_{11}(s^n)$ from (iii), (iv) and (vi) in observations 1. Also, since $\sum_{t=0}^{\infty} \delta^{t+1} s_{n-t}^n > 0$, and thus $q_0(s^n 1) > 0$, $p_{10}(s^n) > p_{11}(s^n)$ from (viii) and (ii). Finally, notice that $\sum_{t=0}^{\infty} \delta^{t+1} s_{n-t}^n$ for these sequences are the partial sums of the alternating geometric series with ratio δ , hence $\sum_{t=0}^{\infty} \delta^{t+1} s_{n-t}^n < 1$, implying $q_0(s^n 0) < 0$ and hence $p_{01}(s^n) > p_{00}(s^n)$ from (vii) and

²⁰Smaller than $(\sqrt{5} - 1)/2 \approx 0.618$, since that's when $\delta + \delta^2 = 1$.

(i). Putting these together means that both of the following orderings must hold:²¹

$$p_{01}(s^n) > p_{10}(s^n) > p_{11}(s^n)$$

$$p_{01}(s^n) > p_{00}(s^n) > p_{11}(s^n)$$

The main conclusion here is that $p_{10}(s^n)$ and $p_{00}(s^n)$ are the intermediate beliefs, $p_{11}(s^n)$ is always seen as the least likely result, and $p_{01}(s^n)$ is the belief with highest probability. This shows that the extended model reinforces the results from the original for these alternating sequences. Beliefs are such that a exact continuation of the pattern of alternation is seen as the most likely result from tossing the coin twice again.

Additional Observations

As mentioned, the same expressions that permit establishing inequalities between outcome-specific beliefs also allow to establish patterns of deviation from the correct uniform belief. The following registers some.²²

Observations 2. *In the extended Rabin and Vayanos model, the following hold for any sequence $s^n \in 2^\omega$:*

- i. $q_0(s^n) < 0 \Rightarrow p_{00}(s^n) < 0.25$ and $p_{01}(s^n) < 0.25 \Rightarrow p_{10}(s^n) > 0.25$ or $p_{11}(s^n) > 0.25$,
- ii. $q_0(s^n) > 0 \Rightarrow p_{10}(s^n) < 0.25$ and $p_{11}(s^n) < 0.25 \Rightarrow p_{01}(s^n) > 0.25$ or $p_{00}(s^n) > 0.25$

5.1.3 Comparing Complexities of Four Future Sequences

Revisit the information-theoretic measure of complexity that was presented in section 2.2. When considering the single next outcome, the comparison was between the

²¹Both $p_{00}(s^n) > p_{10}(s^n)$ and $p_{10}(s^n) > p_{00}(s^n)$ are possible. For $s^n = 01$, for example, $\alpha = 0.5$ leads to $p_{00}(s^n) > p_{10}(s^n)$ for any δ , while $\delta = 0.5$ leads to $p_{00}(s^n) > p_{10}(s^n)$ for any α .

²²Again, from algebraic manipulations one obtains the equivalences for any $R \in \{=, >, <\}$: $p_{00} R 0.25 \Leftrightarrow -q_0(s^n) R q(s^n0)(1 + q(s^n))$ and $p_{01} R 0.25 \Leftrightarrow -q_0(s^n)q(s^n0) R q(s^n0) - q(s^n)$. (i) then follows, establishing also (ii) by symmetry.

complexities of the two sequences that would result from the next outcome being one or another. That is, the comparison was made between $H_k(s^n0)$ and $H_k(s^n1)$, for $k = 0, 1, 2$. If, for example, $H_k(s^n0) > H_k(s^n1)$, that would mean that a next outcome of 0 would make a resulting sequence that is more complex, more random, more representative of the uniform distribution, in order k , than an outcome of 1. Therefore, it should be expected that 0 is more likely than 1. This was captured by variables of the type $\Delta_0 H_k(s^n) = H_k(s^n0) - H_k(s^n1)$.

The same logic applies to the case of the two next outcomes. The only difference being that now the comparison is made between four possible resulting sequences. That is, the objects considered are

$$H_k(s^n00), H_k(s^n01), H_k(s^n10), H_k(s^n11), k = 0, 1, 2.$$

Two immediate, related concerns appear with this extension. The first is that there is no longer a single dimension of comparison between possible future sequences, because they now vary along two instead of only one new addition to the sequence. This makes a one-dimensional variable like $\Delta_0 H_k(s^n)$ not be immediately available. The second is that with additional outcomes, it additionally opens possibilities for tradeoffs between between one order of entropy and another. These make it harder to establish what is the prediction that the entropy approaches make.

To approach such concerns and partially help address them, consider imposing a partial order among the possible outcomes for each sequence s^n . Let $\nu, \nu' \in \{00, 01, 10, 11\}$ be two possible outcomes. Say that ν **entropy-dominates** ν' for sequence s^n if $H_k(s^n\nu) \geq H_k(s^n\nu')$ for all $k \in \{0, 1, 2\}$ and $H_k(s^n\nu) > H_k(s^n\nu')$ for some $k \in \{0, 1, 2\}$. Say that an outcome ν is **entropy-dominated** if there's another outcome ν' that entropy dominates it.

Given a sequence s^n , the ideal scenario is to have an outcome ν which entropy-dominates the other three. In such a scenario, the prediction is unambiguous. The second-best scenario is to eliminate some possible outcomes, potentially to two. For this purpose, one can *eliminate* any entropy-dominated outcome and consider the remaining ones.

Among sequences of length 2, 4 and 6, a process of eliminations like this results in the following. The number of sequences which have a single best outcome which entropy-dominates all others are: 2 out of 4 for length 2, 10 out of 16 for length 4, and 30 out of 64 for length 6.²³ Only 10 sequences of length 6,²⁴ have three outcomes which are not entropy-dominated. In all others, the elimination results in at most two outcomes remaining.

Streak Sequences

The complexity approach likewise confirms an ambiguity regarding the prediction for the streak sequences given by all 1s or all 0s. What might be surprising is that the ambiguity is not, say, between 11 and 10 for a sequence like $s^6 = 000000$. Looking at sequences of 2, 4, 6 and 8 0s, one sees that the outcome 11 produces the highest zero-order entropy, but it's in fact the outcome 01, first continuing the streak and then breaking it, which produces the highest entropy of higher orders. The prediction then hinges on the relative weight given for zero or higher order.

Sequence	Future Sequence	H_0	H_1	H_2
00	0000	0.0	0.0	0.0
	0001	0.811278	0.688721	0.0
	0010	0.811278	0.5	0.0
	0011	1.0	0.5	0.0

Sequence	Future Sequence	H_0	H_1	H_2
0000	000000	0.0	0.0	0.0
	000001	0.650022	0.601606	0.540852
	000010	0.650022	0.540852	0.459147
	000011	0.918295	0.540852	0.459147

²³Or 32 out of 64 for length 6, as for the sequences 011100 and 100011 the outcomes 01 and 10 result in equal entropies and both entropy-dominate the remaining two.

²⁴Which are 001101, 011000, 011011, 011110, 011111 and their symmetrical mirrors.

Sequence	Future Sequence	H_0	H_1	H_2
000000	00000000	0.0	0.0	0.0
	00000001	0.543564	0.517714	0.487517
	00000010	0.543564	0.487517	0.451205
	00000011	0.811278	0.487517	0.451205

Sequence	Future Sequence	H_0	H_1	H_2
00000000	0000000000	0.0	0.0	0.0
	0000000001	0.468996	0.452933	0.434852
	0000000010	0.468996	0.434852	0.414171
	0000000011	0.721928	0.434852	0.414171

Alternating Sequences

When considering alternating sequences, the outcomes which exactly matches the pattern are always entropy-dominated. The outcomes which ‘switch’ the patterns are never entropy-dominated and always produce higher zero order entropy. For lengths 2 and 4, this is the outcome which entropy-dominates all others. For lengths 6 and 8, there is a tradeoff between zero and higher order entropies. For length 6, the other outcome which is not entropy-dominated is either 00 or 11. For length 8, both 00 and 11 actually generate the same entropies.

Sequence	Future Sequence	H_0	H_1	H_2
01	0100	0.811278	0.5	0.0
	0101	1.0	0.0	0.0
	0110	1.0	0.5	0.0
	0111	0.811278	0.0	0.0

Sequence	Future Sequence	H_0	H_1	H_2
0101	010100	0.918295	0.459147	0.333333
	010101	1.0	0.0	0.0
	010110	1.0	0.459147	0.333333
	010111	0.918295	0.459147	0.333333

Sequence	Future Sequence	H_0	H_1	H_2
010101	01010100	0.954434	0.405639	0.344361
	01010101	1.0	0.0	0.0
	01010110	1.0	0.405639	0.344361
	01010111	0.954434	0.5	0.344361

Sequence	Future Sequence	H_0	H_1	H_2
01010101	0101010100	0.970951	0.360964	0.324511
	0101010101	1.0	0.0	0.0
	0101010110	1.0	0.360964	0.324511
	0101010111	0.970951	0.485475	0.324511

5.2 Experiment

The second experiment heavily built on top of the original one. Subjects were likewise presented with sequences of outcomes generated by a computerized version of a fair coin. The sequences were all randomized, with lengths varying from zero to eight outcomes, in increments of two, with the order of their appearance also random. The difference is that subjects were asked about the next two tosses of the coin, instead of simply the next one. The main concerns remained the same, which included not falling into the pitfalls that have constrained previous experiments in the literature. Mostly, the goal is once again to produce rich data that allows the analysis to consider individual sequences. Moreover, there remained the objective of exploiting the structure provided by asking subjects for both choice and belief responses while requiring them to be consistent.

The extension to asking subjects about the next two tosses significantly complicated the space of possible responses. Instead of being asked to choose one among two options and providing one number which was enough to characterize their belief distributions, now subjects faced four outcomes for their choice and their beliefs were three-dimensional. Furthermore, if subjects were simply asked to pick the most likely outcome out of four and were then asked to provide beliefs consistent with this choice, this would do little to actually discipline their belief responses. The task implemented a tiered response structure to address these issues, to be explained next.

5.2.1 Task

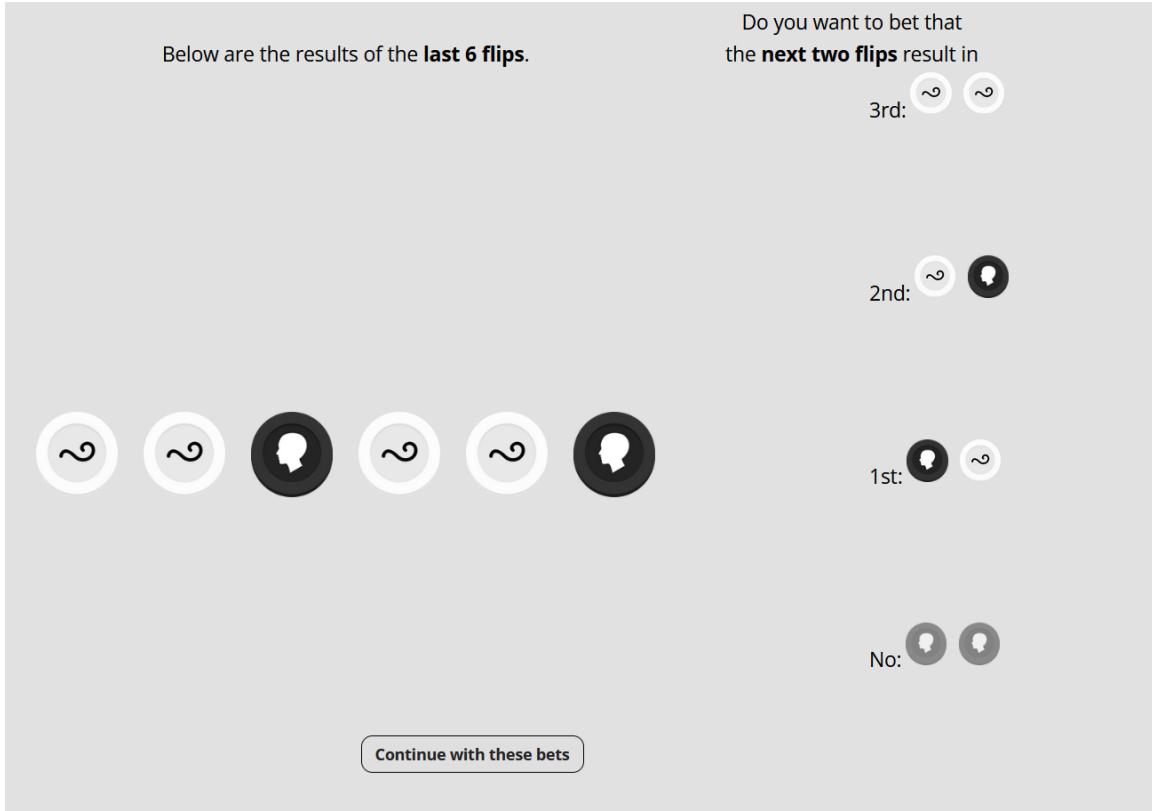


Figure 16: Typical task screen for a given round in the choice stage. Subjects are shown sequences of heads and tails and are asked for the next two results. The center-left of the screen shows the sequence of outcomes and remains the same throughout the round. On the right side of the screen subjects clicked on the possible outcomes to place three bets in the order that they preferred. The bottom button was used to submit their bets and proceed.

Figure 16 shows what a typical task screen looked like for the choice stage. In the center-left of the screen, identically to the first experiment, a sequence of heads and tails was presented to the subject.²⁵ Subjects had to place *three bets* on the possible outcomes, which were ordered as a first bet, a second bet and a third bet. One outcome was left without a bet. Subjects placed these bets via selecting or unselecting the outcomes, which they did by clicking on them. The order was initially

²⁵A tail outcome was coded as 0, a head outcome was coded as 1. So, for example, the sequence appearing in figures 16 and 17 is 001001.

derived by the order in which they clicked, but they could click again on a selected outcome to unselect it and replace the bet. Once they were happy with their bets, they clicked the button at the bottom of the screen to proceed. The two next tosses were then generated and compared to the submitted bets. If they matched the first bet, the largest bonus is awarded. If they matched the second bet, an intermediate bonus is awarded. If they matched the third bet, the smallest bonus was awarded. And if they matched the outcome without a bet, no bonus was awarded.

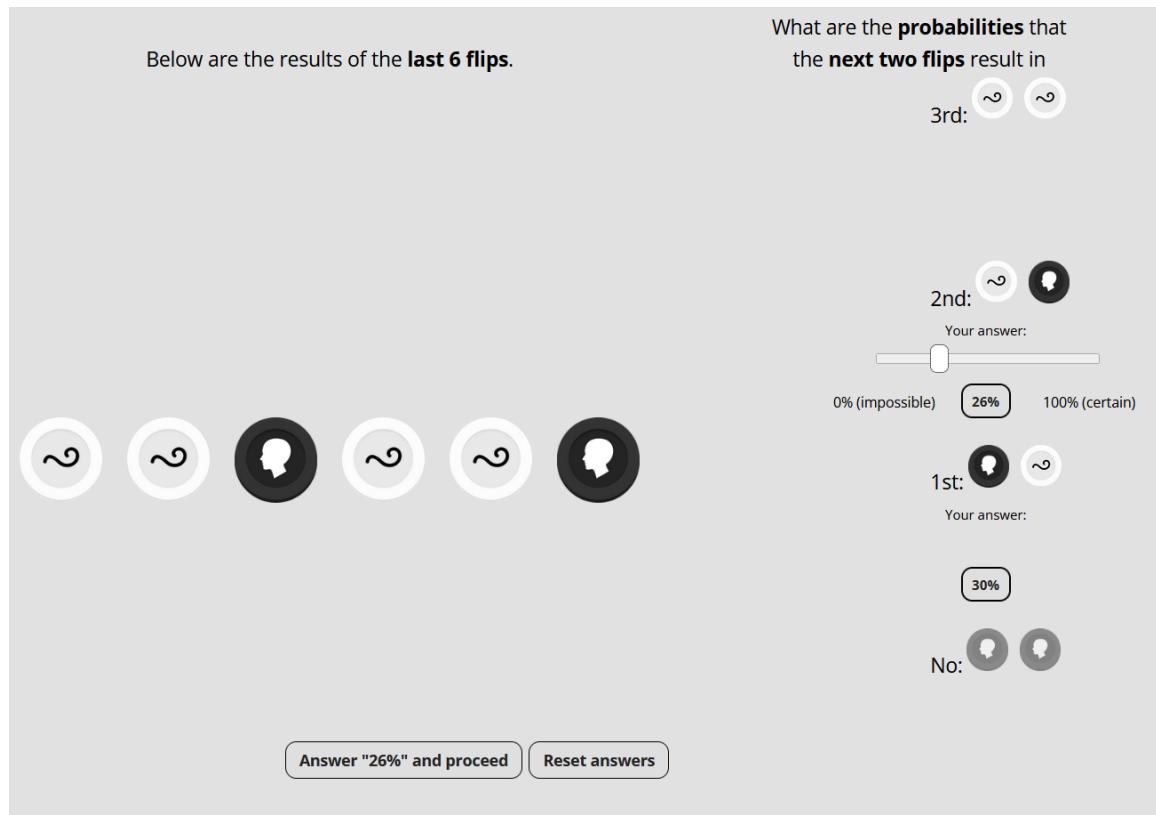


Figure 17: Typical task screen for a given round in the probabilities stage. The center-left of the screen remains as in the choice stage. On the right side of the screen subjects clicked on the possible outcomes to place three bets in the order that they preferred. The left-bottom button was used to submit their bets and proceed. The right-bottom button was used to reset their probability answers and start over from the first bet.

Subjects then proceeded to submit probabilities for each outcome. Figure 17 shows a typical probability screen, continuing from figure 16. They submitted probabilities

in the order of their bets. That is, they first submitted the probability for the outcome on which they placed their first bet, then did it for the second bet, and finally the third. The outcome without a bet received the remaining probability out of 100. When giving the probability for the third bet outcome, the remaining probability being allocated to the no bet outcome was also shown. Probabilities were submitted by interacting with a slider. The sliders were programmed to guarantee the consistency of the probabilities with their bets and the mathematical validity of their total answers.²⁶ If subjects were unsatisfied with their probability answers before submitting the final answer, they could reset them to start over from the first bet.²⁷ Figure 18 shows a possible progression of probability answers.

²⁶More specifically, each slider had upper and lower bounds which subjects could not drag them past. Let $sliderno = 1, 2, 3$ be the number of the slider, $totalprob$ be the total probability already allocated for the previous bets and $previousprob$ be the probability allocated to the immediate predecessor slider. The lower bounds were given by $\lceil(100 - totalprob)/(5 - sliderno)\rceil$. The upper bound for the first slider was 100 and for the second and third, they were $\min\{previousprob, 100 - totalprob\}$.

²⁷The bets were already fixed at this stage and could not be changed, however.

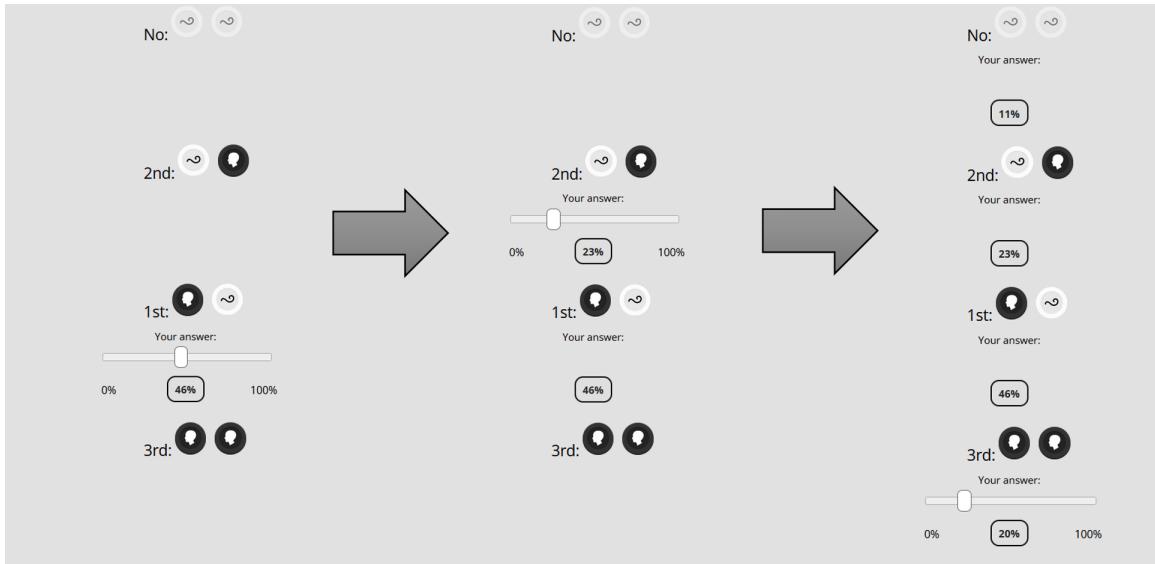


Figure 18: A potential progression of probability answers given that the first bet was put on head and tail, the second on tail and head, and the third on head and head, with tail and tail being left with no bet. The first slider puts probability 46% on head and tail. The second slider puts probability 23% on tail and head. The third slider puts probability 20% on head and tail, which then leaves 11% to be automatically assigned to tail and tail.

Subjects received no feedback.

5.2.2 Structure

After consenting to participate, subjects were presented with instruction screens explaining the nature of the task to them and how their answers would affect their bonus payment.²⁸ Subjects were explained that the computer was producing the sequences they would be seeing in a way that exactly mirrored a real fair coin. The instructions first presented an explanation of the choice response, which was followed by 6 ‘warm up’ questions, in which subjects only had to choose the next outcome. Out of these 6 questions, two were chosen and if their first bet matched the actual result, \$0.60 was added to their bonus payment, \$0.40 was added if their second bet

²⁸Other than two instead of one practice question being selected for payment and the basis of payment being adapted as explained in the previous subsection, this structure is identical to the original one. It is mostly reproduced here for the reader’s convenience.

matched the result, and \$0.20 was added if it was their third bet that matched.²⁹ Then the probability response was introduced and explained to them. At the end of the instructions, there was one practice question with both choice and probability responses just like a normal round. The bonus structure for choice and probabilities was explained to them and is detailed in the next subsection.

Subjects went through 50 rounds in which they encountered, in random order, 2 sequences of length 0³⁰, 4 sequences of length 2, 12 sequences of length 4, 20 sequences of length 6 and 12 sequences of length 8. In each of these sequences, each outcome was randomized, i.e. it was equally likely to be a head or a tail. After these 50 sequences, subjects encountered 4 more sequences, one of each of the previous lengths and in random order. Each one of these sequences was composed of either all heads or all tails, each case being equally likely. For compensation purposes, these last 4 sequences were treated just like the previous 50, and this was explained to subjects.³¹

Before finishing the experiment, subjects were asked the same non-incentivized multi-choice followup question asking for their reasoning in responding to the full run sequences, just like in the original experiment. They were presented with 6 possible answers. See again figure 2.

5.2.3 Implementation

407 subjects were recruited in the Prolific online platform. Using this platform allowed for easy recruiting this large number of participants from a broad population. Subjects were restricted to being located in the USA and being fluent in English.

Throughout the instructions, subjects were asked 4 simple comprehension questions about the functioning of the task and what the responses meant. In accordance

²⁹This effectively amounted to an extra \$0.30 fixed payment in expectation regardless of their actual choices and was only meant to fix their understanding of the task before introducing belief responses.

³⁰The question language was similar but slightly changed from the one seen in figures 16 and 17 for length 0. Subjects were asked ‘Do you want to bet on the **first two flips** resulting in’ and ‘What is the **probability** that the **first two flips** result in’.

³¹Given that these last four sequences were not randomized in the same way as the previous 50, the language was adjusted to account for this. Subjects were asked what would their answers “have been” if they “had seen” these sequences.

with Prolific’s policies on comprehension questions, subjects who answered the same question incorrectly twice are not permitted to continue. This happened to 18 subjects.

The fixed payment for finishing the experiment is \$6.00. The bonus payment varies from \$0.00 to \$14.00. A potential bonus of \$1.20 came from the warm-up questions. Four questions out of the 54 were randomly selected for the remaining bonus. The choice response was worth a maximum of \$1.00 bonus. If their first bet matched the actual result, it generated a \$1.00, \$0.67 was generated if their second bet matched the result, and \$0.33 was generated if it was their third bet that matched.³² The probability response was worth a \$2.20 bonus, calculated according to the Binarized Scoring Rule (Hossain and Okui 2013). Following recent results from Danz et al. (2022), subjects were not directly given explicit mathematical formulas for this calculation, but were told that their bonus payment was maximized by truthfully reporting their best guesses for the correct probabilities. Mathematical formulas were provided at the end of the experiment.

The experiment was coded in JavaScript using custom plug-ins for the jsPsych library (de Leeuw (2015)).

5.3 Results

5.3.1 Revisiting Previous Results

The first few subsections revisit the main results from the original experiment, with the main object of analysis being the reduced beliefs about the probabilities of the next outcome being either 0 or 1:

$$p_{reduced}(s^n) = (p_0(s^n), p_1(s^n)) = (p_{00}(s^n) + p_{01}(s^n), p_{10}(s^n) + p_{11}(s^n)).$$

Figure 19 shows these reduced beliefs and choice for full run sequences for all subjects. The choices exhibit a gambler’s fallacy effect for length up to 6, while beliefs actually present the opposite effect for lengths 4 to 8. However, as might be expected, this is the result of subject heterogeneity, to be explored next.

³²So expected bonus for the choice stage was always \$0.50 as in the original experiment.

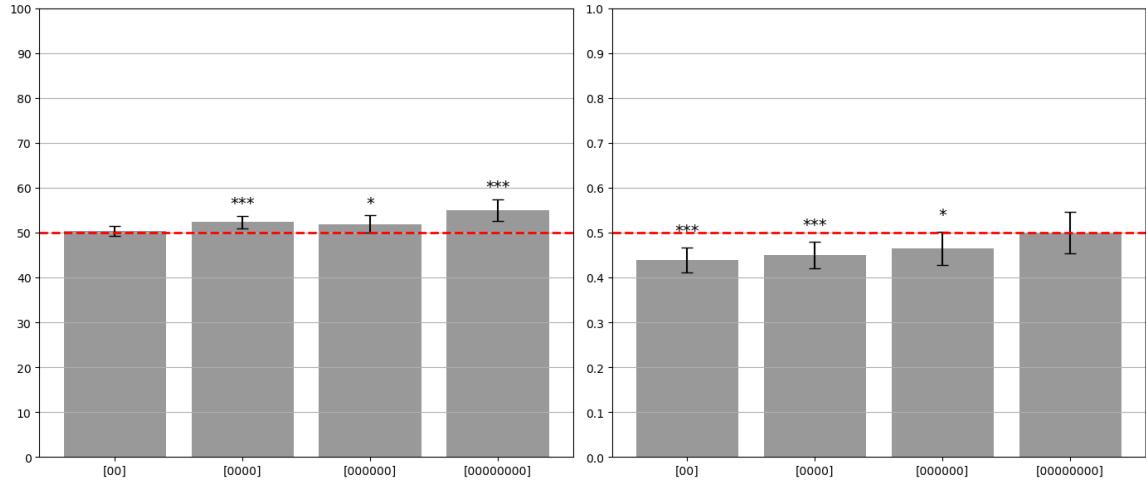


Figure 19: Mean reduced probability (left) and choice frequency (right) responses for the outcome that continues the full run. All subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

5.3.2 Subject Heterogeneity

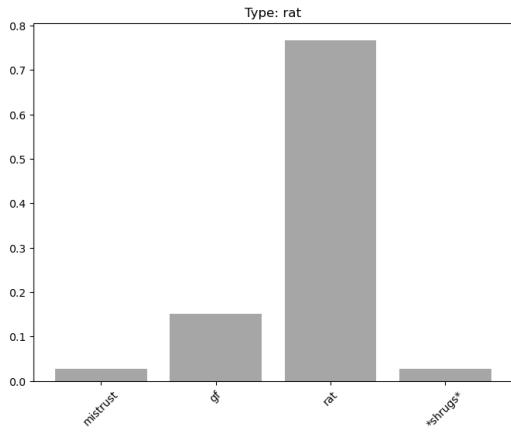
As in the original experiment, subjects were classified into the four types based on their answers on the final four rounds. Table 6 shows the distribution of types for this new experiment. Compared to the original experiment, the proportion of *gambler’s fallacy* subjects remains similar (28.61% before vs. 29.48% now), but this group becomes the largest, since *rational* type subjects decrease from the most common type at 29.62% to the least common at 17.94%. *Hot hands* subjects go from the least common type at 19.52% to the second most common type at 26.54%. *Both* subjects go from 22.22% to 26.04%.

Type	Number	Proportion
gf	120	29.48%
hh	108	26.54%
both	106	26.04%
rat	73	17.94%

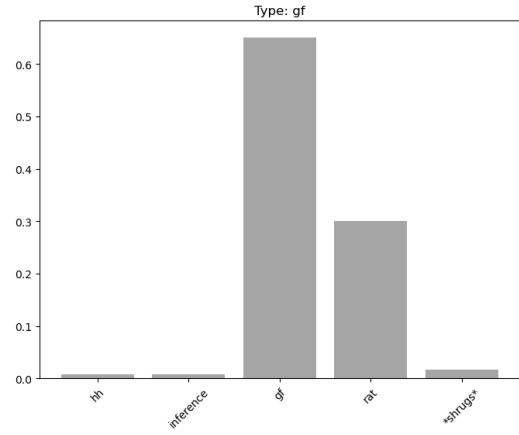
Table 6: Distribution of subject types based on their reduced probability answers.

Figure 20 shows the distribution of answers for the followup question asked at the

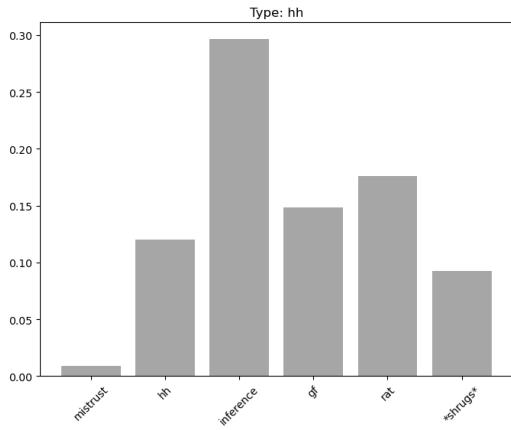
very end of the experiment, separated by the types of subjects. For ‘rat’ (figure 20a) and ‘gf’ (figure 20b) type subjects, their answers continue to be strongly aligned with their behavior. For ‘hh’ (figure 20c) type subjects, the ‘inference’ answer is now the most common, with the ‘hh’ answer falling behind ‘gf’ and ‘rat’ answers, these last two contradicting their actual behavior. Another thing to note is the strong reduction in ‘mistrust’ answers, potentially because this new task asking for the next two coin tosses looks less suspiciously obvious than it might have looked for some subjects who were asked about just the next toss in the original experiment.



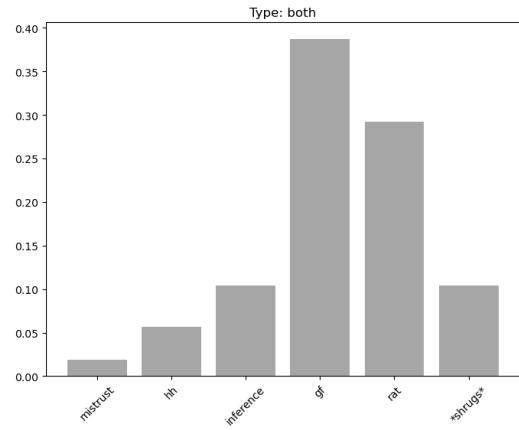
(a) Distribution of follow-up responses for ‘*rational*’ type subjects.



(b) Distribution of follow-up responses for ‘*gambler’s fallacy*’ type subjects.



(c) Distribution of follow-up responses for ‘*hot hands*’ type subjects.



(d) Distribution of follow-up responses for ‘*both*’ type subjects.

Figure 20: Distribution of follow-up responses for different subject types.

Figures 21, 22, 23 and 24 show results for full run sequences discriminated by type. These mostly reproduced the respective figures from the original experiment, including ‘rational’ subjects exhibiting a gambler’s fallacy effect in their choice, although a weaker effect in beliefs for ‘gambler’s fallacy’ subjects is noted, and this is confirmed in the next step of the analysis.

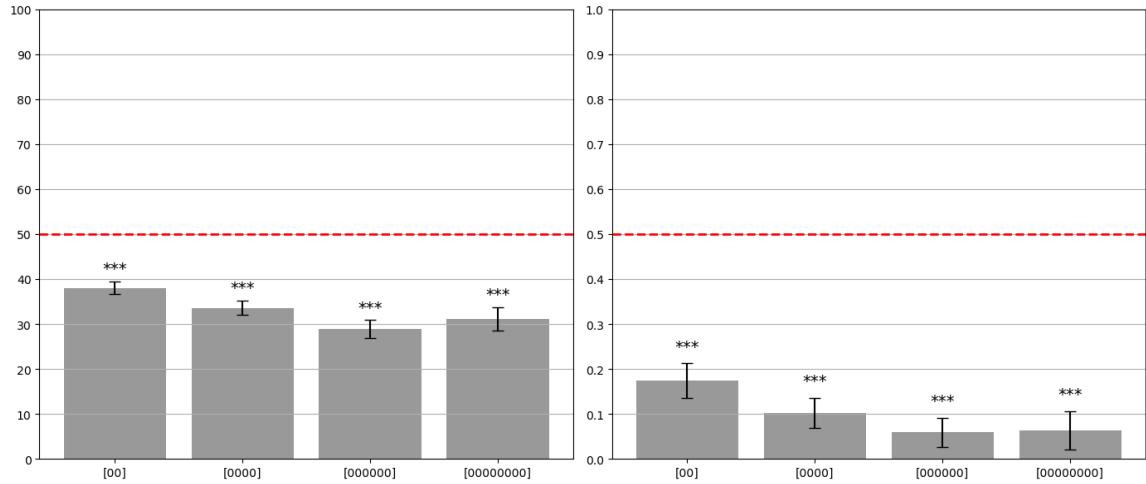


Figure 21: Mean reduced probability (left) and choice frequency (right) responses for the outcome that continues the full run. ‘Gambler’s fallacy’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

5.3.3 Reduced Rabin and Vayanos

Table 7 reports the estimated parameters from a nonlinear least squares estimation of the Rabin and Vayanos model, as done for the original experiment. A significant reversal parameter $\alpha = 0.176853$ is found (a weaker reversal effect than the $\alpha = 0.224$ found for the original experiment), and a significant recency parameter $\delta = 0.737$ is found (with a stronger recency effect than the $\delta = 0.795$ than in the original experiment). These parameters are used to provide the estimate represented as red dots in subsequent figures.

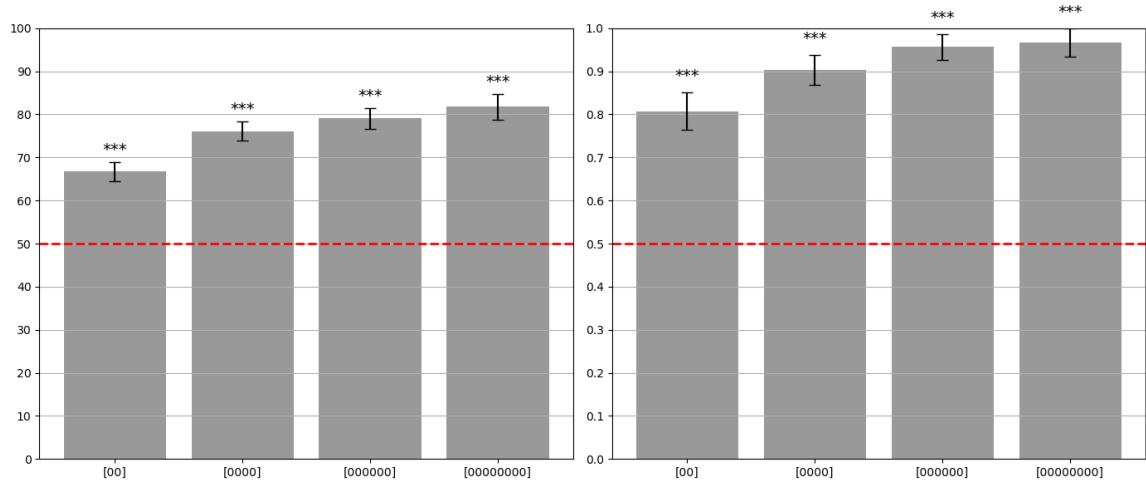


Figure 22: Mean reduced probability (left) and choice frequency (right) responses for the outcome that continues the full run. ‘Hot hands’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

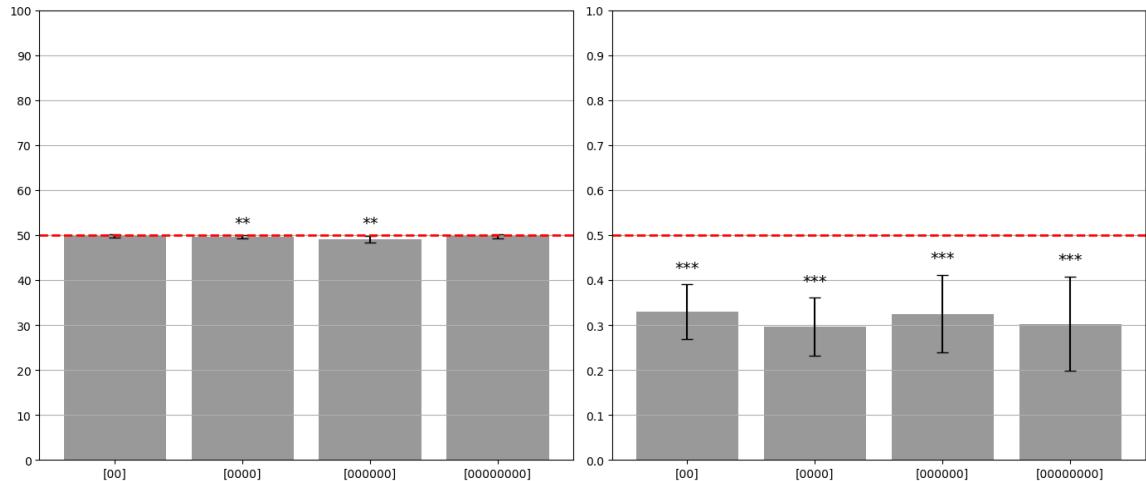


Figure 23: Mean reduced probability (left) and choice frequency (right) responses for the outcome that continues the full run. ‘Rational’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

5.3.4 Alternating Sequences and Monotonicity

Figure 25 presents results for the alternating sequences, with similar results to the original experiment. ‘Gambler’s fallacy’ subjects believe in a continuation of the

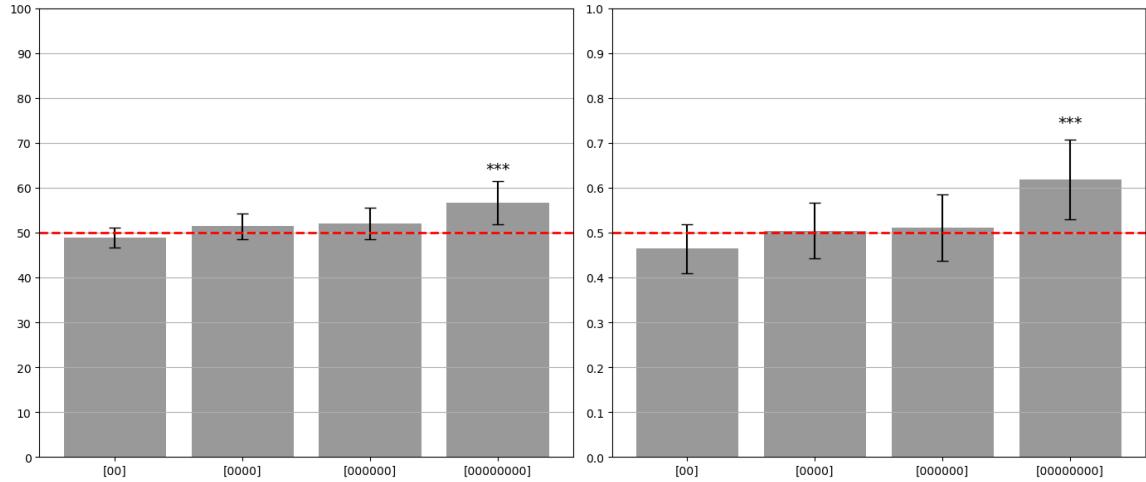


Figure 24: Mean reduced probability (left) and choice frequency (right) responses for the outcome that continues the full run. ‘Both’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

Parameter	Estimate
α	0.176853*** (0.013663)
δ	0.737553*** (0.022895)
β_{length}	-0.000820 (0.001171)
β_{round}	0.000391 (0.000199)
N	6480

Table 7: Results for nonlinear least squares estimation of the Rabin and Vayanos (2010) model for ‘gf’ type subjects using their reduced beliefs. Standard errors clustered at the subject level. * significant at 10% level, ** significant at 5% level, *** significant at 1% level.

alternating pattern for the next outcome, rather than a break. Figures 26 and 27 present results for sequences of almost-full runs, with only one different outcome of varying position. These are again in line with the monotonicity found in the original experiment for such sequences, which isolates the recency bias that ‘gambler’s fallacy’

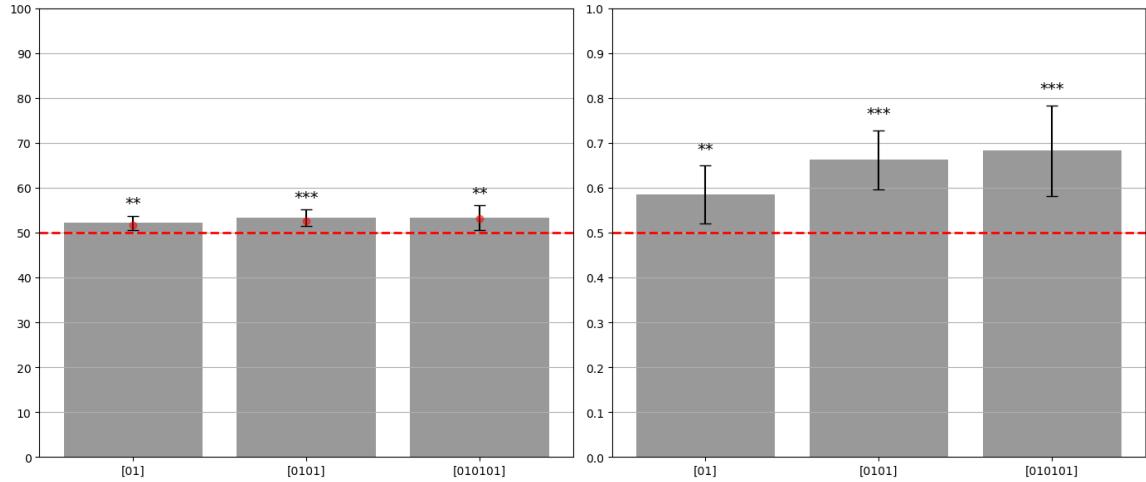


Figure 25: Mean reduced probability (left) and choice frequency (right) responses for the outcome that continues the alternating pattern. ‘Gambler’s fallacy’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

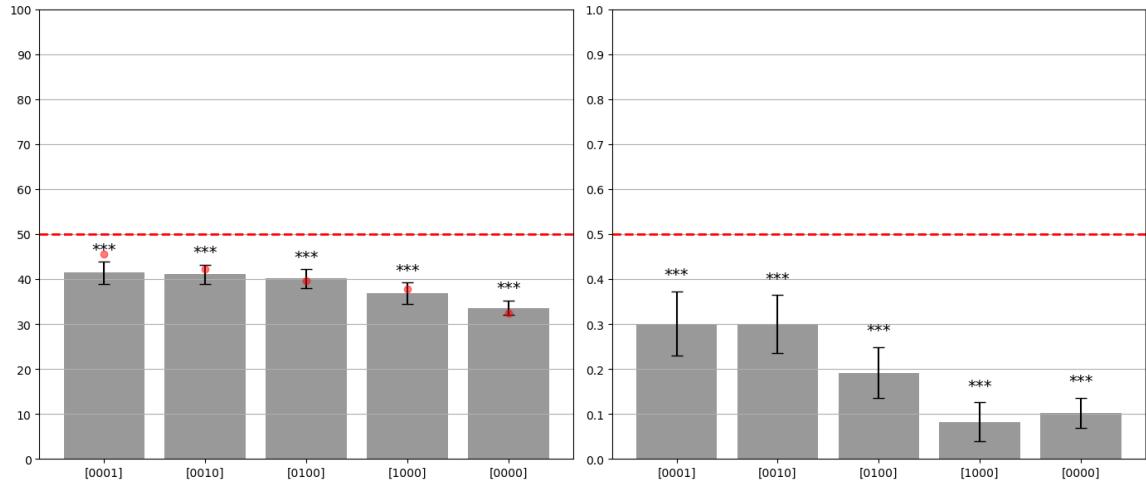


Figure 26: Mean reduced probability (left) and choice frequency (right) responses for length 4 sequences showing monotonicity in responses. ‘Gambler’s fallacy’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

subjects exhibit.

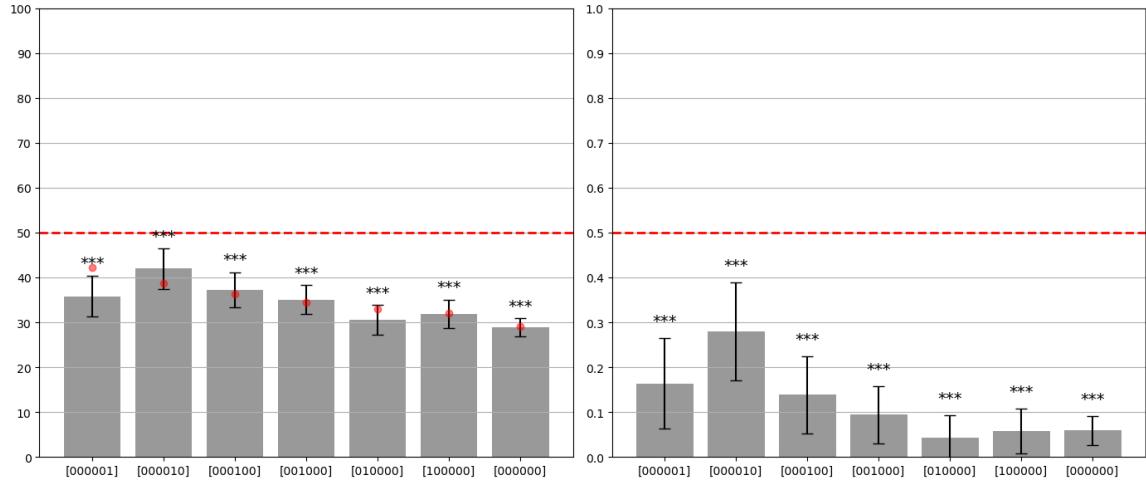


Figure 27: Mean reduced probability (left) and choice frequency (right) responses for length 6 sequences showing monotonicity in responses. ‘Gambler’s fallacy’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

5.3.5 Multidimensional beliefs and choices

Having confirmed that the main results from the original experiment hold, the analysis now proceeds to look at full multidimensional beliefs and choices. When doing so, it carries over the typing of subjects obtained in the previous subsection. The reason for proceeding as such is that, when considering the next outcome, there’s a classic definition of the gambler’s fallacy that has been universal in the literature for decades. Because of the absence of any standard definition of the gambler’s fallacy for predicting the next two outcomes, any method for classifying a subjects as ‘gf’ based on their three-dimensional beliefs will be questionable. Using the typing from the previous subsection preserves the classical definition: someone suffers from the gambler’s fallacy if they believe a tail to be more likely after several heads and vice-versa. What this fallacy implies about beliefs over the next two tosses is then left open to be explored.

5.3.6 Multidimensional Basics

Figure 28 shows mean beliefs and choices for 0-length sequences for all subjects. This figure introduces the new format used for multidimensional beliefs and choices. On the left panel, full beliefs over the four possible outcomes, 00, 01, 10, 11 are presented. They are color-coded, respectively, as white, light gray, dark gray and black. The red dotted line now marks the correct answer of 25%, and means are tested against this value. Red dots, when they appear, mark the prediction given by the extended recency-reversal model presented in subsection 3.2.2, using the parameters obtained with reduced beliefs ($\alpha = 0.176$ and $\delta = 0.737$). The right panel presents the frequencies with which each outcome receives the first bet. That is, the frequency in which it is believed to be the most likely outcome.

Beliefs and choices in figure 28 exhibit a form of gambler's fallacy. The balanced outcomes, with one 0 and one 1, are believed to be more likely than the repeated outcomes of 00 and 11.

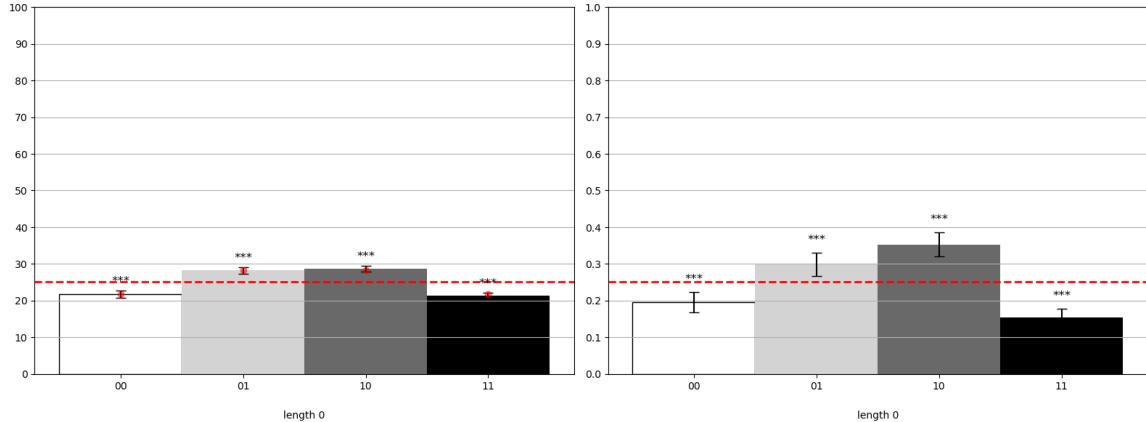


Figure 28: Mean multidimensional probability (left) and choice frequency (right) responses when encountering sequences of length 0. All subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 25 (left) or 0.25 (right): * for 10%, ** for 5%, *** for 1%.

Figure 29 presents multidimensional beliefs and choices for all subjects for full runs. The absence of much in the way of a clear patterns once again motivates

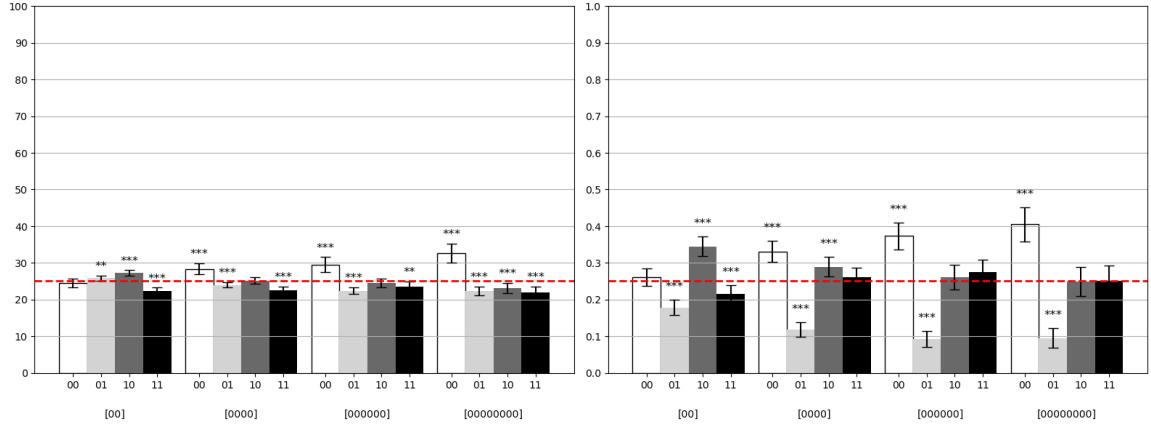


Figure 29: Mean multidimensional probability (left) and choice frequency (right) responses for full runs. All subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 25 (left) or 0.25 (right): * for 10%, ** for 5%, *** for 1%.

looking at versions of this figure for each type of subject.

5.3.7 Subject Heterogeneity

Figure 30 presents multidimensional beliefs and choices for full runs for ‘rat’ subjects. These subjects exhibit again a stark difference in their beliefs and choices. For beliefs, their answers are again, on average, very close to correct. But they exhibit a type of gambler’s fallacy effect in their choices. Incidentally, this effect, in which they put their first bet most often on the outcome that first breaks the run and then returns to it, is actually different from the effect presented by the ‘gf’ subjects themselves.

Figure 31 presents multidimensional beliefs and choices for full runs for ‘hh’ subjects. These subjects present a strong consistency in their belief responses across run length, with a fixed ordering over outcomes, complemented by a monotonicity in their beliefs over the exact double outcome that extends the run.

Figure 32 presents multidimensional beliefs and choices for full runs for ‘both’ subjects. This figure presents the first time, across both experiments, in which a possibility is actually manifested with statistical significance. These subjects switch, over the outcome that extends the run, from a gambler’s fallacy-like belief on the

shortest run to a hot hands-like one for the longest run.

Finally, figure 33 presents multidimensional beliefs and choices for full runs for ‘gf’ subjects. Here the switch case that was encountered in the theoretical section is present. For short runs, the believed-most-likely outcome is that which first breaks the run but does not repeat. For long runs, the repeated outcome that breaks the run is believed most likely. This is inconsistent with the value of δ estimated using the reduced beliefs in the previous subsection. As the red dots indicate, the estimated value would imply a consistent ordering over the outcomes in beliefs.

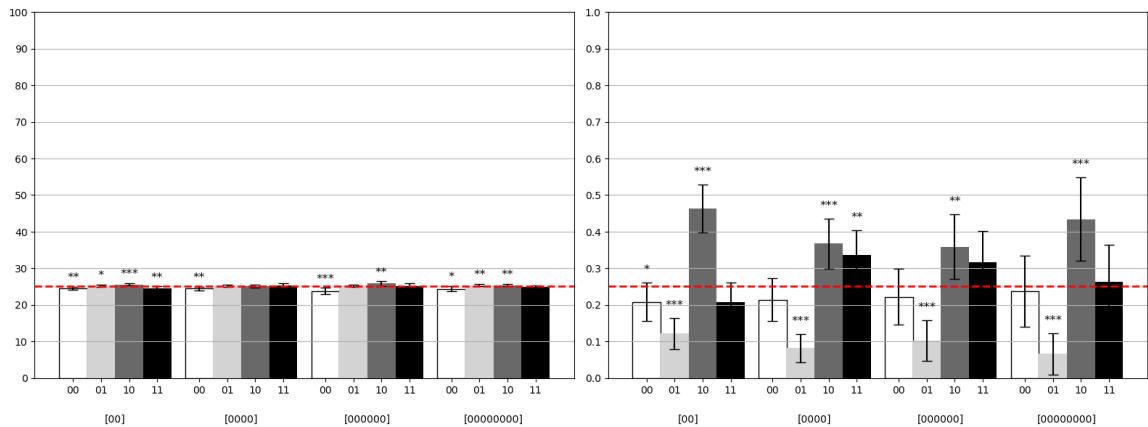


Figure 30: Mean multidimensional probability (left) and choice frequency (right) responses for full runs. ‘Rational’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 25 (left) or 0.25 (right): * for 10%, ** for 5%, *** for 1%.

5.3.8 Alternating Sequences

Figure 34 presents the results, for ‘gf’ subjects, for the alternating sequences. Beliefs strongly support the extended version of the Rabin and Vayanos model, and they are reinforced by the choice results, while contradicting any prediction based on the breaking of patterns. Subjects believe that the most likely two outcomes are those that exactly reproduce the ongoing alternating pattern.

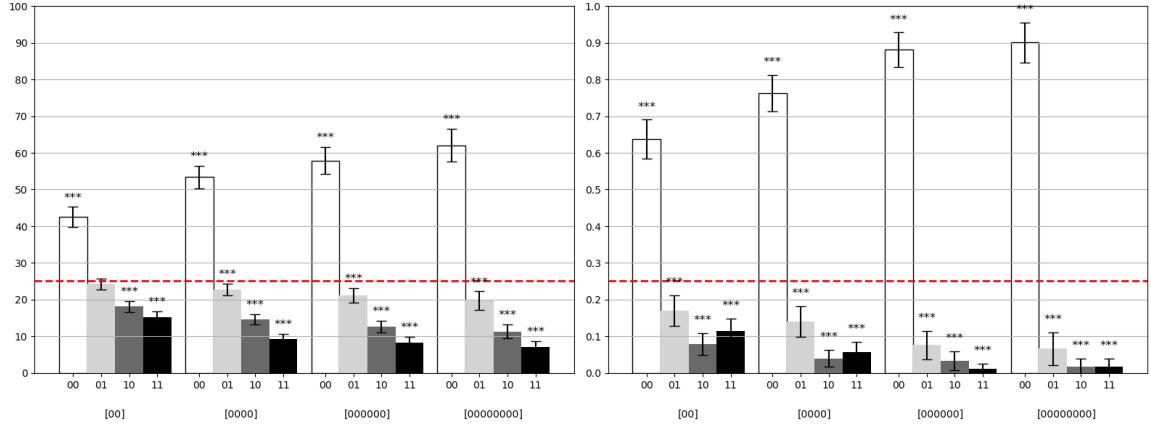


Figure 31: Mean multidimensional probability (left) and choice frequency (right) responses for full runs. ‘Hot hands’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 25 (left) or 0.25 (right): * for 10%, ** for 5%, *** for 1%.

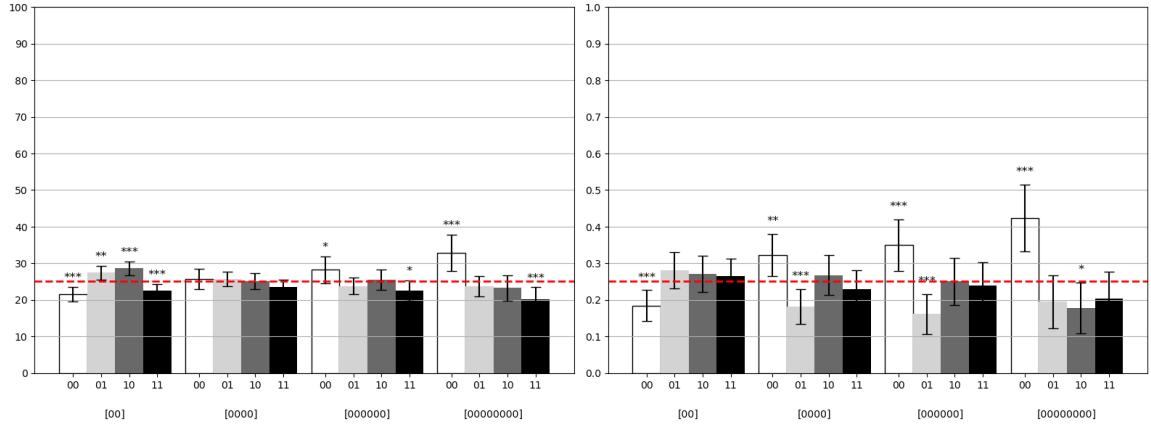


Figure 32: Mean multidimensional probability (left) and choice frequency (right) responses for full runs. ‘Both’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 25 (left) or 0.25 (right): * for 10%, ** for 5%, *** for 1%.

6 Discussion and Conclusion

In this paper I have compared the performance of two theoretical models of the gambler’s fallacy: a formalized version of the representativeness heuristic using information theory, and the recency-weighted reversal model of Rabin and Vayanos (2010). Although both models perform well when taking an aggregate look at the

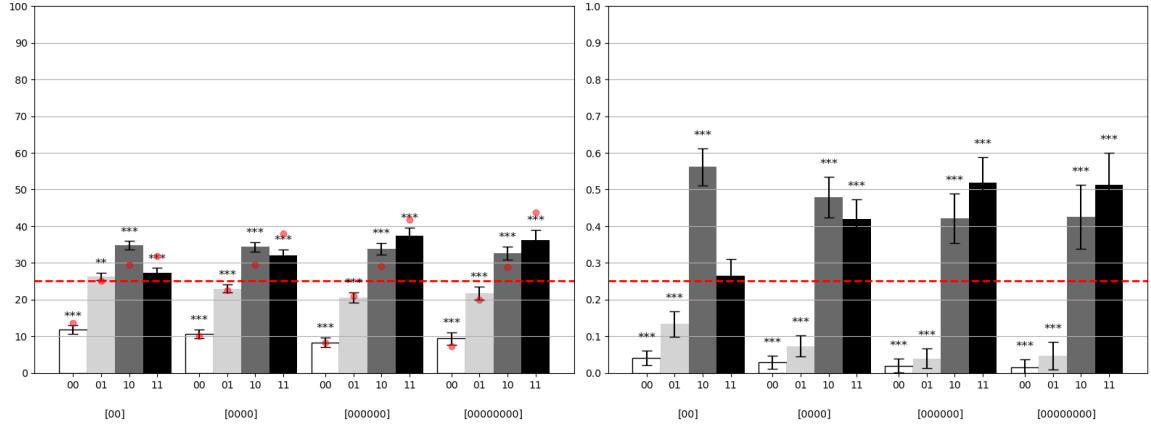


Figure 33: Mean multidimensional probability (left) and choice frequency (right) responses for full runs. ‘Gambler’s fallacy’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 25 (left) or 0.25 (right): * for 10%, ** for 5%, *** for 1%.

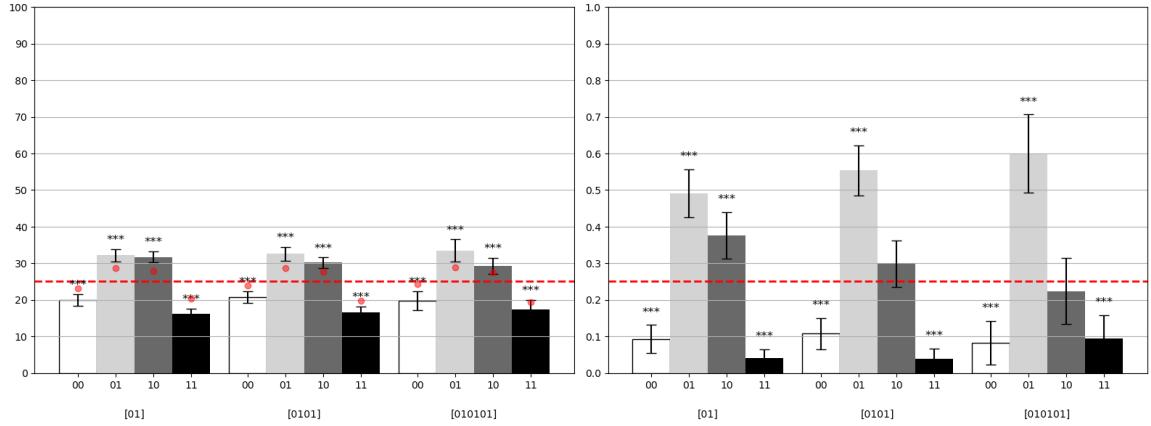


Figure 34: Mean multidimensional probability (left) and choice frequency (right) responses for alternating sequences. ‘Gambler’s fallacy’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 25 (left) or 0.25 (right): * for 10%, ** for 5%, *** for 1%.

data, more detailed analysis at the level of individual sequence reveals that the latter has a superior performance.

The main component of this superior performance comes from the clear recency bias that subjects exhibit in their behavior. Subjects’ beliefs were more responsive to the outcomes appearing more to the front of the sequences, whereas those further back had a lesser effect. Previous experiments, which did not collect belief responses

for general sequences, could not definitively confirm this bias in the context of the gambler’s fallacy. This behavior adds to the long standing literature on recency bias going back to Sternberg (1966). Afrouzi et al. (2023) provide, along with their own model and experimental results, an overview of the literature. Both in their model and in more recent work they survey,³³ the recency bias is modeled as resulting from costly cognitive processing of information in working memory formation. This extends to cases in which no recall is actually required, as in this paper’s experiment, since there’s an important component of attention control to working memory (Unsworth and Spillers 2010). Future work on the gambler’s fallacy will have to engage with this literature on the recency bias, potentially explicitly taking into account its foundations in costly cognition.

Natural extensions of the current work would be to expand the space of either possible outcomes or responses, or the underlying data generating process more generally. One such extension which is planned for this project is to run a very similar experiment with a fair coin, but ask subjects for the probabilities of the next two outcomes (that is, a belief $q \in \Delta(\{00, 01, 10, 11\})$). It’s an open question which of the results from the current experiment will hold in this new version. In particular, what will subjects say is the most likely of the four outcomes when facing the alternating sequences *head-tail-head-tail*.... If the Rabin and Vayanos model extends to this setup, subjects should still think *head-tail* is most likely next. It might be, however, that in this case it becomes more evident that they are continuing the first order pattern of the sequence and therefore judge another outcome, which increases complexity, as more likely to happen. Other extensions would be to use a biased, instead of a fair coin, or a process with more than two outcomes to go beyond the binary case.

Finally, the complexity measure of higher order empirical entropy which has received attention here might find other uses as a more general measure of information complexity. It might find applications in measuring costly information processing and its relations to information preferences.

³³See also their online appendix.

References

- Afrouzi, Hassan, Spencer Y Kwon, Augustin Landier, Yueran Ma, and David Thesmar**, “Overreaction in Expectations: Evidence and Theory,” *The Quarterly Journal of Economics*, March 2023, 138 (3), 1713–1764.
- Bar-Hillel, Maya**, “What features make samples seem representative?,” *Journal of Experimental Psychology: Human Perception and Performance*, August 1980, 6 (3), 578–589.
- , “Representativeness and fallacies of probability judgment,” *Acta Psychologica*, March 1984, 55 (2), 91–107.
- Benjamin, Daniel, Don Moore, and Matthew Rabin**, *Biased Beliefs About Random Samples: Evidence from Two Integrated Experiments* October 2017.
- Dale, Andrew I. and Pierre-Simon Laplace**, *Philosophical Essay on Probabilities*, Springer New York, 1825/1995.
- Danz, David, Lise Vesterlund, and Alistair J. Wilson**, “Belief Elicitation and Behavioral Incentive Compatibility,” *American Economic Review*, September 2022, 112 (9), 2851–2883.
- de Leeuw, J. R.**, “jsPsych: A JavaScript library for creating behavioral experiments in a web browser,” *Behavior Research Methods*, 2015, 47 (2), 1–12.
- Ferragina, Paolo and Giovanni Manzini**, “Indexing compressed text,” *Journal of the ACM*, July 2005, 52 (4), 552–581.
- Gagie, Travis**, “Large alphabets and incompressibility,” *Information Processing Letters*, September 2006, 99 (6), 246–251.
- Grether, David M.**, “Bayes Rule as a Descriptive Model: The Representativeness Heuristic,” *The Quarterly Journal of Economics*, November 1980, 95 (3), 537.
- Hossain, T. and R. Okui**, “The Binarized Scoring Rule,” *The Review of Economic Studies*, January 2013, 80 (3), 984–1001.
- Kahneman, Daniel and Amos Tversky**, “Subjective probability: A judgment of representativeness,” *Cognitive Psychology*, July 1972, 3 (3), 430–454.
- Manzini, Giovanni**, “An analysis of the Burrows—Wheeler transform,” *Journal of the ACM*, May 2001, 48 (3), 407–430.
- Oskarsson, An T., Leaf Van Boven, Gary H. McClelland, and Reid Hastie**, “What’s next? Judging sequences of binary events.”, *Psychological Bulletin*, 2009, 135 (2), 262–285.
- Rabin, Matthew**, “Inference by Believers in the Law of Small Numbers,” *The Quarterly Journal of Economics*, August 2002, 117 (3), 775–816.
- and Dimitri Vayanos, “The Gambler’s and Hot-Hand Fallacies: Theory and Applications,” *Review of Economic Studies*, April 2010, 77 (2), 730–778.
- Shannon, C. E.**, “A Mathematical Theory of Communication,” *Bell System Technical Journal*, July 1948, 27 (3), 379–423.
- Skinner, B. F.**, “The processes involved in the repeated guessing of alternatives,” *Journal of Experimental Psychology*, June 1942, 30 (6), 495–503.

- Sternberg, Saul**, “High-Speed Scanning in Human Memory,” *Science*, August 1966, 153 (3736), 652–654.
- Tenenbaum, Joshua B. and Thomas Griffiths**, “The rational basis of representativeness,” in J. Moore and K. Stenning, eds., *Proceedings of the 23rd Annual Conference of the Cognitive Science Society*, Erlbaum Mahwah, NJ 2001, pp. 1036–1041.
- Tversky, Amos and Daniel Kahneman**, “Judgment under Uncertainty: Heuristics and Biases: Biases in judgments reveal some heuristics of thinking under uncertainty.” *Science*, September 1974, 185 (4157), 1124–1131.
- and —, “Judgments of and by representativeness,” in Daniel Kahneman, Paul Slovic, and Amos Tversky, eds., *Judgment under Uncertainty*, Cambridge University Press, 1982, pp. 84–98.
- Unsworth, Nash and Gregory J. Spillers**, “Working memory capacity: Attention control, secondary memory, or both? A direct test of the dual-component model,” *Journal of Memory and Language*, May 2010, 62 (4), 392–406.

A Individual Sequences

This appendix provides the complete look at individual sequences of lengths 2, 4 and 6.

Figures are similar to the ones shown in subsection 4.1 of the main text. Probability responses are presented on the left and choice responses on the right. Except when noted, these responses are for the outcome that results in lower entropy. The red dots shown on the left panels mark the point-prediction of the Rabin and Vayanos model for the sequence when using the parameters estimated in subsection 4.3 ($(\alpha, \delta) = (0.224, 0.795)$).

Sequences are organized according to their length and also on the orders for which empirical entropy matters. For example, for some sequences, outcomes 0 and 1 only make a difference at the level of empirical entropy of order 1. That is, they are sequences s such that $\Delta H_1(s) > 0$ but $\Delta H_0(s) = 0$ and $\Delta H_2(s) = 0$. All results are for ‘gf’ type subjects only.

A.1 Length 0

Throughout the experiment, subjects encountered twice sequences of length 0. That is, the screen had no outcomes shown, and subjects were asked to bet on the first outcome being a head or a tail, as well as give the probabilities that the first outcome is head or tail. Figure A.1 shows responses when subjects encountered this length 0 sequences for the first time, for the second time, and aggregating across both. Responses are for outcome 0. The slight bias towards 1 (head) probably reflects the fact that in the task screen, this outcome was lower and therefore closer to the confirmation button at the bottom. (See figure 1.) Importantly, the second time they encountered such sequences, subjects were more likely to be careful in giving the correct 50-50 answer.

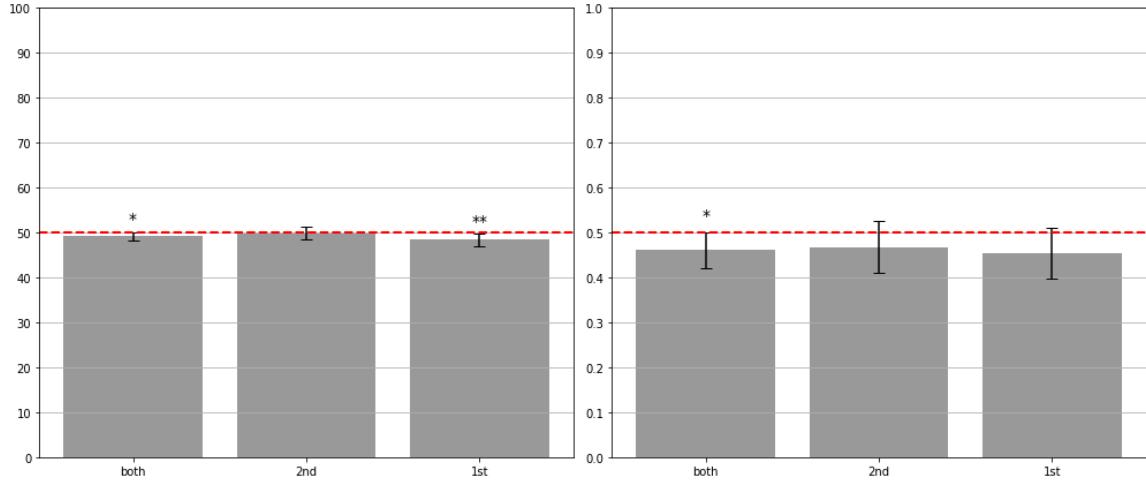


Figure A.1: Mean probability (left) and choice frequency (right) responses for 0 when encountering sequences of length 0 for the first and second times. All subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

A.2 Length 2

Figure A.2 shows responses for outcome 0 for the length 2 sequences. Although not significantly different from 50%, the mean probability responses for the sequences 10 and 01 already serve as a preview from the effect that recency has on subjects' responses.

A.3 Length 6: $\Delta H_0 > 0$, $\Delta H_1 > 0$, $\Delta H_2 > 0$

Figure A.3 shows responses for low entropy outcome for length 6 sequences that have $\Delta H_0 > 0$, $\Delta H_1 > 0$, $\Delta H_2 > 0$. For these sequences, both theories perform well. Subjects generally believe that low entropy outcomes are less likely to come next.

A.4 Length 6: $\Delta H_0 > 0$, $\Delta H_1 > 0$, $\Delta H_2 = 0$

Figure A.4 shows responses for low entropy outcome for length 6 sequences that have $\Delta H_0 > 0$, $\Delta H_1 > 0$, $\Delta H_2 = 0$. These sequences continue the theme from the previous

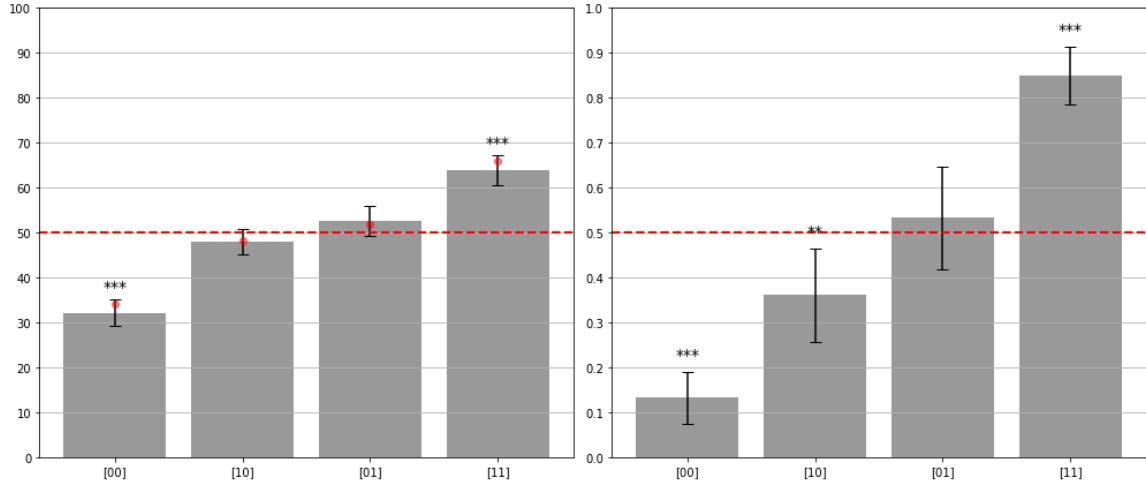


Figure A.2: Mean probability (left) and choice frequency (right) responses for outcome 0 for all length 2 sequences. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

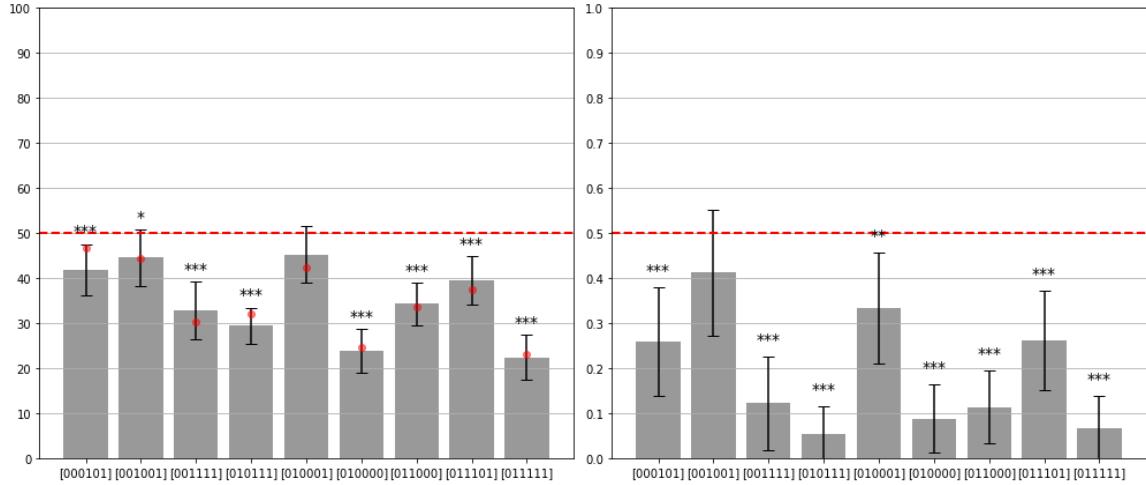


Figure A.3: Mean probability (left) and choice frequency (right) responses for lower entropy outcome for length 6 sequences with $\Delta H_0 > 0$, $\Delta H_1 > 0$, $\Delta H_2 > 0$. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

figure, with subjects thinking that low entropy outcomes are less likely.

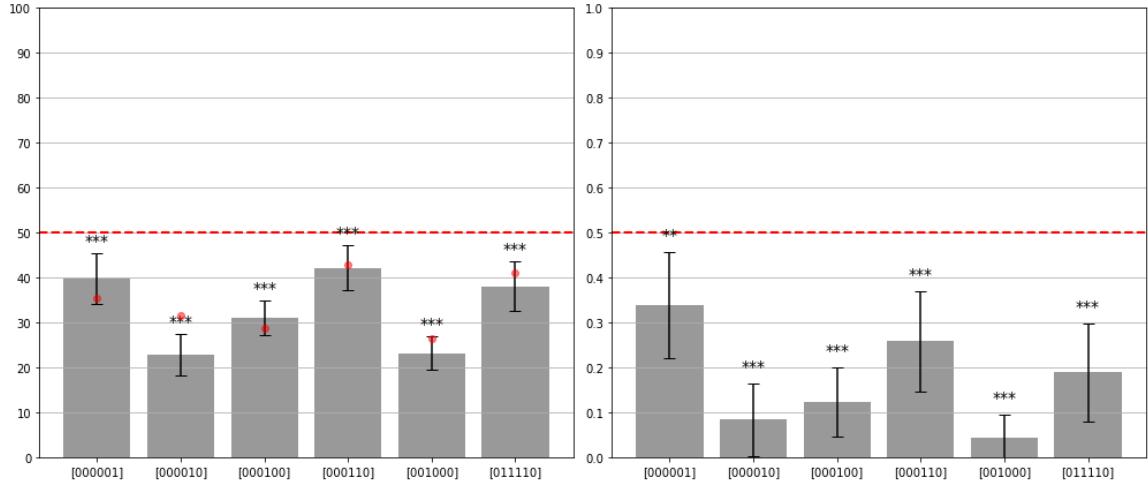


Figure A.4: Mean probability (left) and choice frequency (right) responses for lower entropy outcome for length 6 sequences with $\Delta H_0 > 0$, $\Delta H_1 > 0$, $\Delta H_2 = 0$. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

A.5 Length 4: $\Delta H_0 = 0$, $\Delta H_1 > 0$, $\Delta H_2 = 0$

Figure A.5 shows responses for low entropy outcome for length 4 sequences that have $\Delta H_0 = 0$, $\Delta H_1 > 0$, $\Delta H_2 = 0$. It is among these sequences that some of the most significant violations of the entropy theory happen. For sequences 0101, 1010, 0110 and 1001, subjects believe that the low entropy outcome is more likely to come next. While these responses violate the entropy prediction, they are in line with the prediction made by the Rabin and Vayanos model. For example, for the alternating sequences, 0101 and 1010, because of the recency effect, one of the outcomes is further to the front in all instances and so gets overweighted relative to the other one, leading subjects to believe it less likely to come next.

A.6 Length 6: $\Delta H_0 = 0$, $\Delta H_1 > 0$, $\Delta H_2 > 0$

Figure A.6 shows responses for low entropy outcome for length 6 sequences that have $\Delta H_0 = 0$, $\Delta H_1 > 0$, $\Delta H_2 > 0$. Although at a lower significance level, the alternating sequence here further confirms the same effect as seen in the previous figure. Subjects

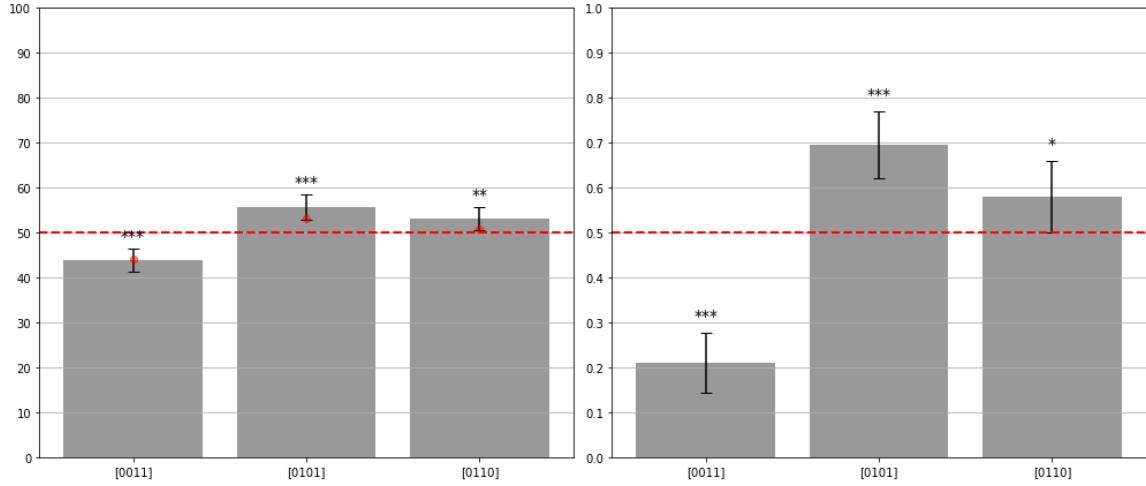


Figure A.5: Mean probability (left) and choice frequency (right) responses for lower entropy outcome for length 4 sequences with $\Delta H_0 = 0$, $\Delta H_1 > 0$, $\Delta H_2 = 0$. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

responses go in the opposite direction to that predicted by the entropy model, while being in line with the prediction of the Rabin and Vayanos model.

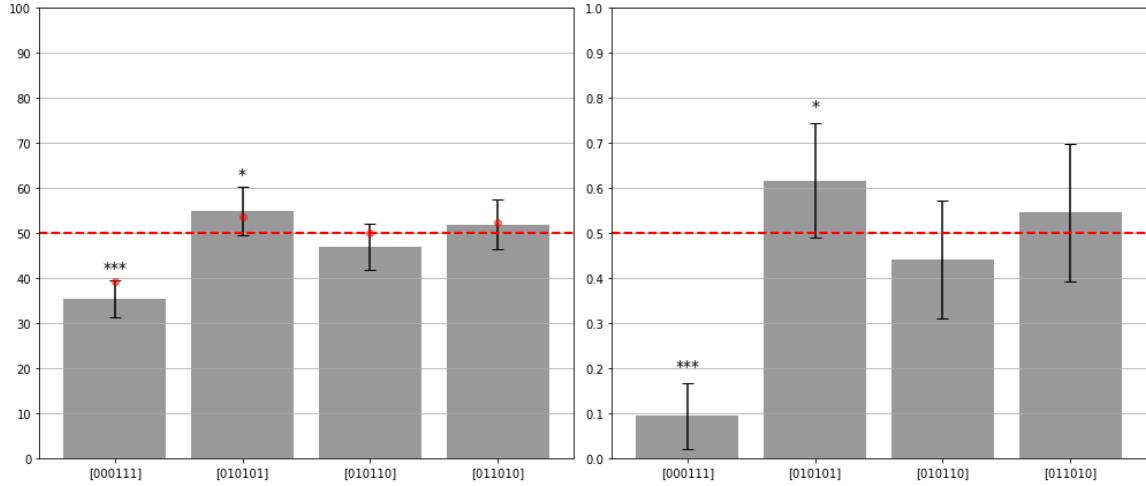


Figure A.6: Mean probability (left) and choice frequency (right) responses for lower entropy outcome for length 6 sequences with $\Delta H_0 = 0$, $\Delta H_1 > 0$, $\Delta H_2 > 0$. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

A.7 Length 6: $\Delta H_0 = 0$, $\Delta H_1 = 0$, $\Delta H_2 > 0$

Figure A.7 shows responses for low entropy outcome for length 6 sequences that have $\Delta H_0 = 0$, $\Delta H_1 = 0$, $\Delta H_2 > 0$. No significant effect is found for these sequences. Disaggregating the two symmetric sequences 011001 and 100110 makes no difference.

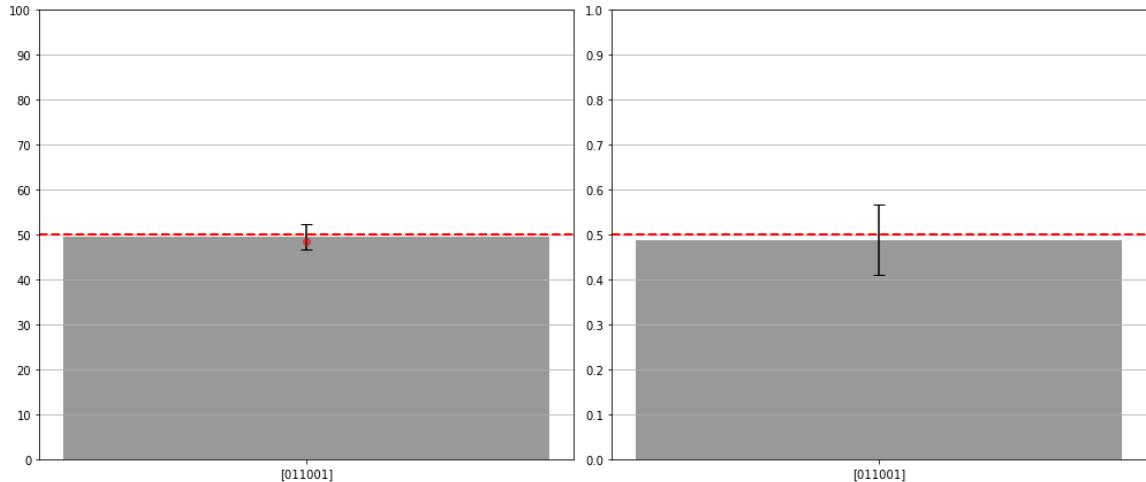


Figure A.7: Mean probability (left) and choice frequency (right) responses for lower entropy outcome for length 6 sequences with $\Delta H_0 = 0$, $\Delta H_1 = 0$, $\Delta H_2 > 0$. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

A.8 Length 6: ΔH_k with different signs

Figure A.8 shows responses for the outcome with low zero order entropy for length 6 sequences that have ΔH_k with different signs. That is, for these sequences, there is a tradeoff between increasing entropy at one order while decreasing it at another. Given the higher impact of lower order entropy, as shown by the estimated parameters in subsection 4.4, it might be expected that it is the outcome that increases zero order entropy that will be thought less likely to happen, which is indeed what is confirmed by these sequences.

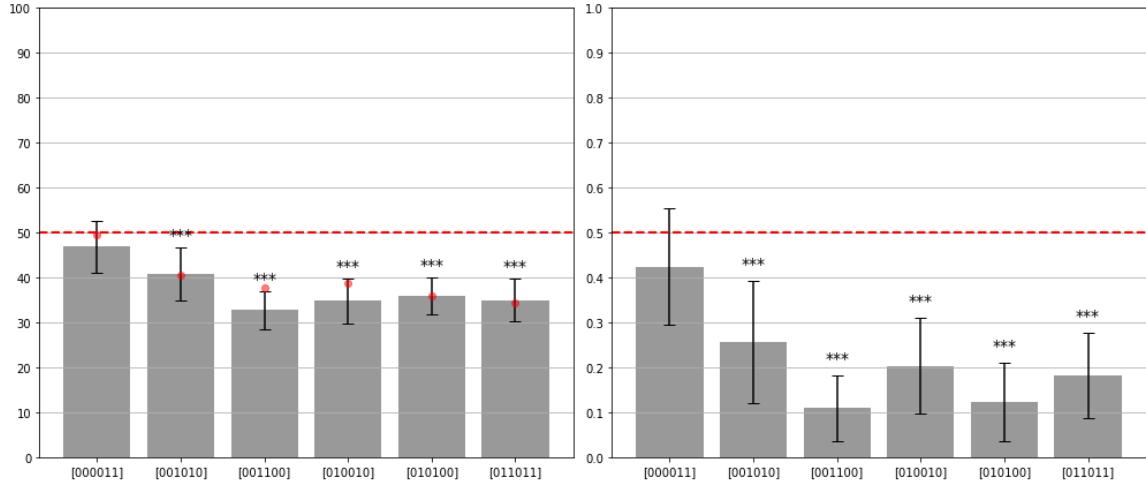


Figure A.8: Mean probability (left) and choice frequency (right) responses for lower zero order entropy outcome for length 6 sequences with ΔH_i having different signs. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

A.9 Length 6: $\Delta H_0 = \Delta H_1 = \Delta H_2 = 0$

Figure A.9 shows responses for outcome 0 for length 6 sequences that have $\Delta H_0 = \Delta H_1 = \Delta H_2 = 0$. That is, these are sequences for which there is no ‘low entropy’ outcome, as both outcomes 0 and 1 result in a sequence with same empirical entropy of orders 0, 1 and 2. These can be thought of sequences that are already highly complex, and there’s little pattern to be matched or reinforced by either outcome. Nevertheless, it’s shown that for most of these sequences, subjects still deviate significantly from the correct belief. Their deviation closely follows the prediction of Rabin and Vayanos. It’s worth noting that this happens with sequences that have the same outcome repeated twice at the end of the sequence, further reinforcing the importance of the recency effect.

A.10 Lengths 4 and 6: recency effects

Figures A.10 and A.11 illustrate in more detail the recency effect for sequences of length 4 and 6, respectively. They show in order sequences that have one (or no)

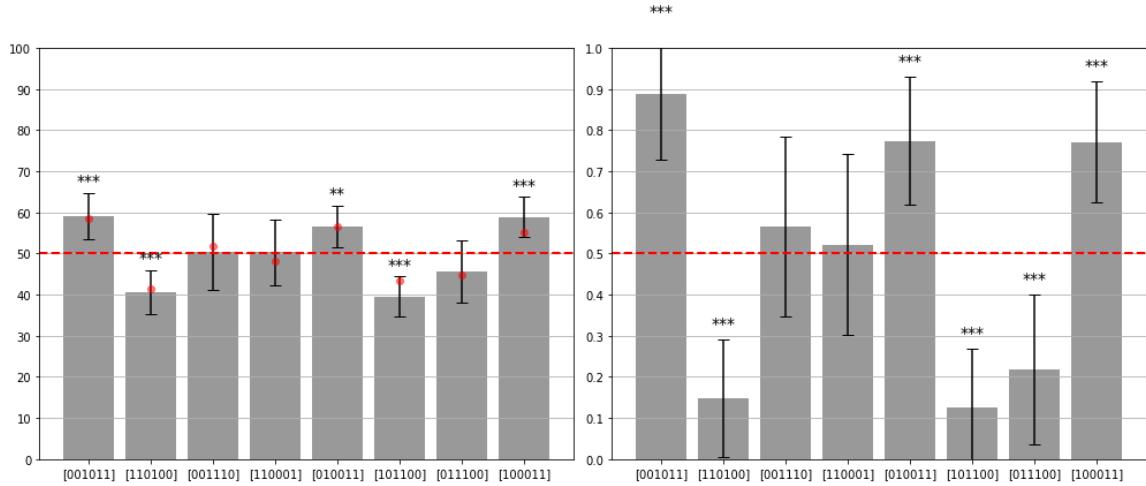


Figure A.9: Mean probability (left) and choice frequency (right) responses for outcome 0 for length 6 sequences with $\Delta H_0 = 0$, $\Delta H_1 = 0$, $\Delta H_2 = 0$. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

outcome different than the rest, but differing by the position in which it appears. These figures show a monotonicity that is created by the recency effect. The closer the different outcome is to the front of the sequence, the larger its impact in countering the effect of the rest of the outcomes on the belief. For example, subjects believe 0 is less likely to be next when facing the sequence 000001 than the sequence 100000. This cannot be explained by empirical entropy, but is in line with the Rabin and Vayanos model.

B Aggregate Regressions

Given the subject heterogeneity explained in section 4.1 of the main text, the data analysis proceeded by treating ‘gf’ and ‘hh’ type subjects separately. This appendix shows results for all subjects, as well as ‘rat’ and ‘both’ subjects.

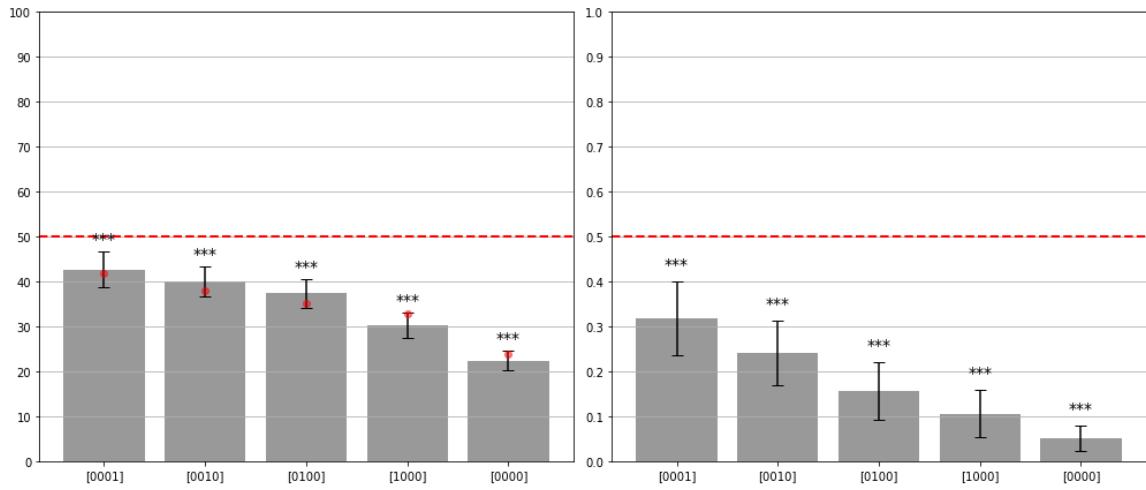


Figure A.10: Mean probability (left) and choice frequency (right) responses for lower entropy outcome, showing recency effect for length 4 sequences. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

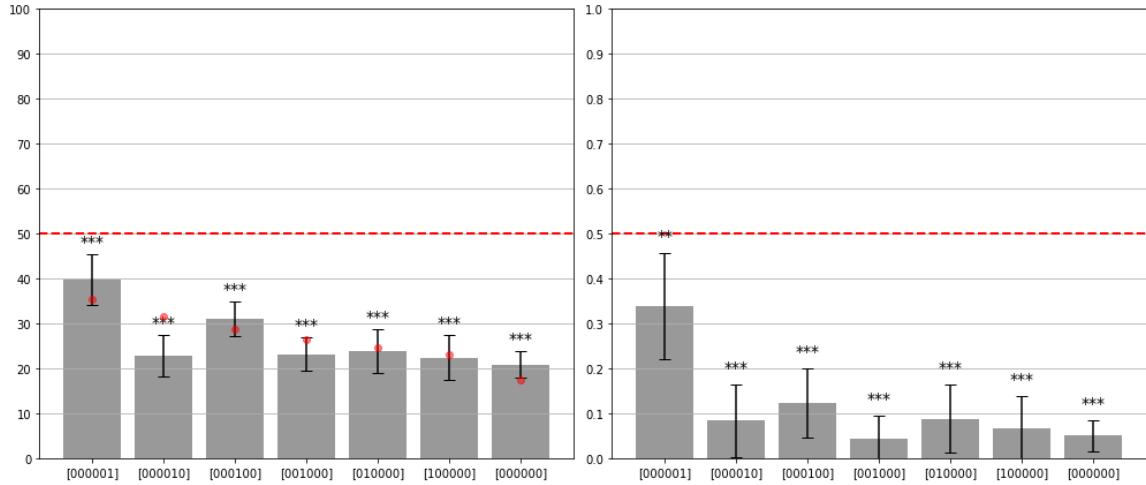


Figure A.11: Mean probability (left) and choice frequency (right) responses for lower entropy outcome, showing recency effect for length 6 sequences. ‘gf’ subjects. Error bars are 95% confidence intervals. Stars record level of significance at which the displayed values are different from 50 (left) or 0.5 (right): * for 10%, ** for 5%, *** for 1%.

B.1 Rabin and Vayanos model

Similarly to section 4.3 in the main text, tables B.1, B.2 and B.3 show fitted parameters for, respectively, all subjects, type ‘both’ subjects and type ‘rat’ subjects. As mentioned in the main text, this functional form is not adequate for ‘hh’ type subjects, and for these subjects the estimated parameters are essentially all 0. See table 5 and figure 15 in the main text for the appropriate analysis.

Parameter	Estimate
α	0.216846*** (0.079945)
δ	0.297822** (0.121728)
β_{length}	0.000004 (0.001054)
β_{rounds}	0.000115 (0.000185)
<i>N</i> = 16038	

Table B.1: Results for nonlinear least squares estimation of the Rabin and Vayanos (2010) model for all subjects according to equation 1. Standard errors clustered at the subject level. * significant at 10% level, ** significant at 5% level, *** significant at 1% level.

Parameter	Estimate
α	0.413009 (0.295782)
δ	0.214402 (0.163174)
β_{length}	0.001061 (0.002720)
β_{rounds}	0.000256 (0.000401)
<i>N</i> = 3564	

Table B.2: Results for nonlinear least squares estimation of the Rabin and Vayanos (2010) model for type ‘both’ subjects according to equation 1. Standard errors clustered at the subject level. * significant at 10% level, ** significant at 5% level, *** significant at 1% level.

Parameter	Estimate
α	0.002218 (0.002965)
δ	1.000000*** (0.206427)
β_{length}	0.001244 (0.000813)
β_{rounds}	-0.000106 (0.000104)
<hr/> $N = 4752$	

Table B.3: Results for nonlinear least squares estimation of the Rabin and Vayanos (2010) model for type ‘rat’ subjects according to equation 1. Standard errors clustered at the subject level. * significant at 10% level, ** significant at 5% level, *** significant at 1% level.

B.2 Empirical Entropy

Similarly to section 4.4 in the main text, tables B.4, B.5 and B.6 show regression results for, respectively, all subjects, type ‘both’ subjects and type ‘rat’ subjects.

B.3 Generalized Rabin and Vayanos

Similarly to section 4.6, table B.7 reports the results of the generalized Rabin and Vayanos model for all subjects.

	q_0
<i>const</i>	0.0027 (0.011)
ΔH_0	0.0400 (0.033)
ΔH_1	0.0731*** (0.023)
ΔH_2	-0.0341 (0.026)
<i>length</i>	-0.0003 (0.002)
<i>rounds</i>	0.00008 (0.000)
<i>R</i> ²	0.003
N	16,038

Table B.4: Regression results for all subjects. q_0 , the dependent variable, is the reported probability of 0 as the next outcome. OLS regression with standard errors clustered at the subject level. * significant at 10% level, ** significant at 5% level, *** significant at 1% level.

	q_0
<i>const</i>	0.0183 (0.031)
ΔH_0	0.0357 (0.047)
ΔH_1	0.1004 (0.062)
ΔH_2	-0.0869 (0.067)
<i>length</i>	-0.0008 (0.004)
<i>rounds</i>	0.00000 (0.001)
<i>R</i> ²	0.002
N	3,564

Table B.5: Regression results for type ‘both’ subjects. q_0 , the dependent variable, is the reported probability of 0 as the next outcome. OLS regression with standard errors clustered at the subject level. * significant at 10% level, ** significant at 5% level, *** significant at 1% level.

	q_0
<i>const</i>	0.0058 (0.006)
ΔH_0	0.0102 (0.011)
ΔH_1	0.0085 (0.010)
ΔH_2	-0.0078 (0.014)
<i>length</i>	0.0006 (0.001)
<i>rounds</i>	-0.0002 (0.000)
<i>R</i> ²	0.001
N	4,752

Table B.6: Regression results for type ‘rat’ subjects. q_0 , the dependent variable, is the reported probability of 0 as the next outcome. OLS regression with standard errors clustered at the subject level. * significant at 10% level, ** significant at 5% level, *** significant at 1% level.

	q_0
<i>const</i>	0.0054 (0.012)
<i>pos₁</i>	0.0650*** (0.009)
<i>pos₂</i>	0.0270*** (0.007)
<i>pos₃</i>	0.0052 (0.007)
<i>pos₄</i>	-0.0117* (0.007)
<i>pos₅</i>	-0.0062 (0.007)
<i>pos₆</i>	-0.0129** (0.006)
<i>pos₇</i>	-0.0170** (0.008)
<i>pos₈</i>	-0.0100 (0.007)
<i>rounds</i>	0.0001 (0.000)
<i>size</i>	-0.0006 (0.002)
<i>R</i> ²	0.026
N	16038

Table B.7: Regression results for all subjects of the generalized Rabin and Vayanos model. q_0 , the dependent variable, is the reported probability of 0 as the next outcome. OLS regression with standard errors clustered at the subject level. * significant at 10% level, ** significant at 5% level, *** significant at 1% level.